

**Lecture Information**

**Course:** CSCI E-89B: Natural Language Processing  
**Lecture:** Lecture 7  
**Topic:** Latent Dirichlet Allocation and Topic Modeling  
**Date:** Fall 2024

**Contents**

## 1 Quiz Review: Autoencoders Revisited

### Overview

This lecture begins with a review of autoencoder concepts before introducing topic modeling—a powerful unsupervised learning technique for discovering hidden themes in document collections.

### 1.1 One-Hot Encoding Disadvantages

#### One-Hot Encoding Problems

One-hot encoding has two major disadvantages:

1. **Sparsity:** The representation contains many zeros
2. **High dimensionality:** As a direct result of sparsity

#### Why sparsity is problematic:

- When computing linear combinations, you perform many operations with zeros
- $0 \times$  something contributes nothing but still requires computation
- Information is represented inefficiently
- This is precisely why we use embeddings instead

### 1.2 Undercomplete Autoencoders

#### Undercomplete Autoencoder

An autoencoder is called **undercomplete** when the dimensionality of the representation (encoding) is **lower** than the dimensionality of the input. This creates a bottleneck that forces compression.

#### Understanding “Undercomplete”

The terminology makes intuitive sense:

- **Complete:** 4 dimensions in  $\rightarrow$  4 dimensions in middle  $\rightarrow$  4 dimensions out
- **Undercomplete:** 4 dimensions in  $\rightarrow$  2 dimensions in middle  $\rightarrow$  4 dimensions out
- **Overcomplete:** 3 dimensions in  $\rightarrow$  7 dimensions in middle  $\rightarrow$  3 dimensions out

If you have a 3-dimensional dataset and map it to 2 dimensions, your representation is “undercomplete” because you cannot fully represent 3D data in 2D without some loss. You’re taking projections rather than keeping complete information.

### 1.3 Stacked Autoencoders

#### Stacked Autoencoder

A **stacked autoencoder** (also called a deep autoencoder) has multiple hidden layers. “Stacked” by definition implies depth—at least 2 hidden layers.

- **Shallow network:** 1 hidden layer only
- **Deep network:** 2 or more hidden layers

#### 1.3.1 Layer-wise Training

##### Important

Layers of stacked autoencoders do not need to be trained together. You can train them one at a time using a “sandwich” approach:

**Phase 1:** Train only input →  $\text{hidden}_1 \rightarrow \text{output}$  (no middle layers)

**Phase 2:** Freeze coefficients, add  $\text{hidden}_2$  between  $\text{hidden}_1$  and output

**Phase 3:** Train middle part while keeping frozen layers fixed

**Phase 4:** Continue stacking more layers

#### 1.3.2 Why Layer-wise Training?

The motivation relates to gradient descent optimization:

#### The Scale Problem

Weights  $W$  for different layers can be on different scales:

- $W_1$  (near input) and  $W_{200}$  (near output) may differ vastly in magnitude
- Cost function level curves become stretched (elliptical rather than circular)
- The gradient  $-\alpha \nabla J$  points away from the true minimum
- Training may diverge or take forever

**Modern solutions to the scale problem:**

1. **Adam optimizer:** Adjusts gradient direction based on local curvature
2. **Batch normalization:** Normalizes signals locally, mitigating weight scale issues
3. **Shortcut connections:** Connect layers to later layers, allowing signals to bypass problematic areas

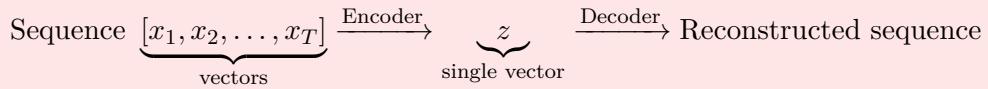
#### Modern Practice

Due to these techniques, the scale problem is largely solved today. People often train deep autoencoders with all layers at once. However, the phase-wise approach remains elegant and worth knowing.

## 1.4 Autoencoders for Sequences

### Important

Autoencoders **can** be used for sequences. A sequence of words (sentence) can be compressed into a single vector (the bottleneck), which has no time component.



### 1.4.1 How Small Should the Bottleneck Be?

This is a fundamental question with no universal answer:

#### Goal Determines Architecture

If your goal is perfect reconstruction:

- Don't squeeze at all!
- Why would you compress if you need to recover everything?
- Just use input = output directly

If your goal is feature extraction for downstream tasks:

- Use the encoder output as input to a classifier
- The optimal bottleneck size depends on classification performance
- Iterate: try 2D, 5D, 10D representations and evaluate which works best

#### Practical Consideration

If you have millions of input features, feeding them directly to a classifier creates millions of parameters. You may want to compress first to:

- Reduce parameter count
- Avoid getting stuck in local minima
- Make optimization tractable

## 2 Introduction to Topic Modeling

### Overview

Topic modeling is an **unsupervised learning** approach for discovering hidden themes in document collections. Unlike clustering, documents can belong to **multiple topics** with different proportions.

## 2.1 Motivation: Why Not Clustering?

### Traditional Clustering

Standard clustering algorithms (K-means, hierarchical) assign each data point to **exactly one cluster**:

- Students belong to “students” cluster
- Retired people belong to “retired” cluster
- Points don’t overlap between clusters

The problem with text documents:

### Documents Are Different

In text, topics naturally **overlap**:

- A news article may start discussing economics, shift to politics, and mention artificial intelligence
- The same document contains multiple themes
- Hard assignment to one cluster loses important information

### Document as Topic Mixture

Consider Document 2 in a corpus:

- Topic 1 (Economics): 60%
- Topic 2 (Politics): 30%
- Topic 3 (AI): 10%

This soft assignment captures the document’s multi-topical nature far better than forcing it into one category.

## 2.2 Topic Modeling Methods

Three main approaches for topic modeling:

Method	Type	Characteristics
LDA	Probabilistic	Assumes stochastic text generation
NMF	Deterministic	Matrix factorization approach
STM	Probabilistic	LDA + covariates (metadata)

### Why Social Scientists Love Topic Models

Topic modeling is popular in political science, government departments, and social sciences:

- Extract information from large text corpora computationally
- Discover biases and patterns

- Correlate topics with covariates (gender, source, etc.)

## 2.3 Latent Topics

### Latent Topics

Topics are called “latent” because:

- We don’t define topics upfront (e.g., “economics”)
- The model discovers topics from data
- We only assign labels **after** processing
- Labels come from examining top-weighted documents/words per topic

### Important

The number of topics is a **hyperparameter** you must specify beforehand. If you say 3 topics, the model will find exactly 3 topics—it won’t tell you the “true” number.

## 3 Maximum Likelihood Estimation

### Overview

Before diving into LDA, we need to understand **maximum likelihood estimation (MLE)**—the framework used to recover model parameters from observed data.

### 3.1 The MLE Framework

#### Maximum Likelihood Estimation

MLE finds parameters that maximize the probability of observing the data we actually observed:

$$\hat{\theta} = \arg \max_{\theta} P(\text{data} | \theta)$$

#### Simple Example: Normal Distribution

Suppose we observe data and create a histogram. We assume data comes from a normal distribution with unknown  $\mu$  and  $\sigma^2$ .

**Step 1:** Try parameters (Model A):  $\mu$  far left of data

- Probability of observing our data given Model A is very small
- Likelihood  $\approx 0$

**Step 2:** Try parameters (Model B):  $\mu$  closer to data center

- Probability is higher
- Likelihood increases

**Step 3:** Try parameters (Model C):  $\mu$  at data mean, appropriate  $\sigma^2$

- Probability is highest
- This is our MLE solution!

**Step 4:** Try parameters (Model D): Overshoot  $\mu$

- Probability decreases again

*Conceptual diagram: Scanning through parameter space to find maximum likelihood*

### Important

MLE is incredibly powerful:

- Works with many parameters (dozens or more)
- Used in time series, LDA, STM, and countless applications
- Very natural approach: find parameters that make observed data most probable

## 4 Latent Dirichlet Allocation (LDA)

### Overview

LDA models each document as a **mixture of topics**, where each topic is a distribution over words. Unlike clustering, LDA provides soft (probabilistic) topic assignments.

### 4.1 The Generative Model

LDA assumes text is generated by the following process:

#### LDA Text Generation Process

For document  $m$ :

**Step 1:** Choose number of words  $N \sim \text{Poisson}(\xi)$

**Step 2:** Choose topic proportions  $\theta_m \sim \text{Dirichlet}(\alpha)$

**Step 3:** For each word  $n = 1, \dots, N$ :

- Choose topic  $z_n \sim \text{Multinomial}(\theta_m)$
- Choose word  $w_n \sim \text{Multinomial}(\beta_{z_n})$

## 4.2 Understanding Each Component

### 4.2.1 Poisson Distribution for Document Length

#### Poisson Distribution

The Poisson distribution models count data:

$$P(N = k) = \frac{\xi^k e^{-\xi}}{k!}$$

where  $\xi$  is both the mean and variance. If documents average 25 words,  $\xi \approx 25$ .

### 4.2.2 Dirichlet Distribution for Topic Proportions

#### Dirichlet Distribution

The Dirichlet distribution produces vectors of probabilities that sum to 1. For  $K$  topics:

$$\theta_m = (\theta_{m,1}, \theta_{m,2}, \dots, \theta_{m,K}) \quad \text{where} \quad \sum_k \theta_{m,k} = 1$$

Example for 3 topics:  $\theta_m = (0.6, 0.3, 0.1)$  means:

- 60% Topic 1 (Economics)
- 30% Topic 2 (Politics)
- 10% Topic 3 (AI)

#### Dirichlet as Generalized Beta

- For 2 topics: Dirichlet reduces to the Beta distribution
- For  $K > 2$  topics: Dirichlet is the natural generalization
- Parameter  $\alpha$  (vector) controls the shape of the distribution

### 4.2.3 Multinomial Distribution for Topic and Word Selection

#### Multinomial Selection

Given probabilities, multinomial sampling selects one category:

- $z_n \sim \text{Multinomial}(\theta_m)$ : Select topic for word  $n$
- $w_n \sim \text{Multinomial}(\beta_{z_n})$ : Select word from chosen topic's vocabulary distribution

#### Concrete Word Generation

Given  $\theta_m = (0.1, 0.2, 0.6)$  for three topics:

1. Draw topic: Most likely Topic 3 (60%), but could be Topic 2 (20% chance)
2. Suppose Topic 2 is selected
3. Draw word from  $\beta_2$ : Each topic has its own word distribution

4. Word “maximization” is selected from Topic 2’s distribution

This is repeated independently for each word position.

### 4.3 Key Assumptions of LDA

#### Bag of Words Assumption

LDA assumes **no word order**:

- Words are generated independently
- “maximization” doesn’t depend on what came before
- Shuffling words in a document doesn’t change its LDA representation
- Generated text wouldn’t be grammatical—that’s not the point!

#### Important

LDA is not for generating readable text. It’s for:

- Discovering what topics a document discusses
- Finding topic proportions for each document
- Identifying words associated with each topic

### 4.4 The EM Algorithm

#### Expectation-Maximization (EM)

LDA uses the EM algorithm because topic assignments  $z_n$  are **latent variables** (not observed):

1. **E-step:** Estimate expected values of latent variables  $z_n$  given current parameters
2. **M-step:** Maximize likelihood given these expected values
3. Iterate until convergence

#### Why EM?

We observe documents (sequences of words) but not:

- Which topic each word came from ( $z_n$ )
- The true topic proportions ( $\theta_m$ )
- The topic-word distributions ( $\beta_k$ )

EM handles this missing information elegantly.

## 5 LDA Implementation in Python

### 5.1 Using scikit-learn

```

1 from sklearn.feature_extraction.text import CountVectorizer
2 from sklearn.decomposition import LatentDirichletAllocation
3
4 # Example documents
5 documents = [
6     "Cats are wonderful pets",
7     "Cats and dogs are popular animals",
8     "Dogs enjoy long walks",
9     "Walks in the park are relaxing",
10    # ... more documents
11 ]
12
13 # Create document-term matrix
14 vectorizer = CountVectorizer(stop_words='english')
15 X = vectorizer.fit_transform(documents)
16
17 # Fit LDA
18 lda = LatentDirichletAllocation(
19     n_components=2,          # Number of topics
20     random_state=42         # For reproducibility
21 )
22 lda.fit(X)
23
24 # Get topic-document distribution
25 doc_topic_dist = lda.transform(X)
26 print(doc_topic_dist)
27 # Output: [[0.05, 0.95], [0.06, 0.94], [0.93, 0.07], ...]

```

Listing 1: LDA with scikit-learn

#### Interpreting Results

Output [0.05, 0.95] means:

- 5% contribution from Topic 1
- 95% contribution from Topic 2

Since this document is mostly Topic 2, and it's about cats, Topic 2 is likely the “cats” topic.

### 5.2 Extracting Top Words per Topic

```

1 def display_topics(model, feature_names, num_top_words=5):
2     for topic_idx, topic in enumerate(model.components_):
3         top_words_idx = topic.argsort()[-num_top_words-1:-1]
4         top_words = [feature_names[i] for i in top_words_idx]
5         print(f"Topic {topic_idx}: {', '.join(top_words)}")
6
7 feature_names = vectorizer.get_feature_names_out()
8 display_topics(lda, feature_names)
9
10 # Output:
11 # Topic 0: dogs, enjoy, long, walks, exploring
12 # Topic 1: cats, purr, love, climb, trees

```

Listing 2: Display top words for each topic

### 5.3 Document-Based Topic Interpretation

#### Important

A better way to understand topics: look at documents with highest topic prevalence, not just top words.

```

1 # For each topic, find documents with highest prevalence
2 for topic_idx in range(n_topics):
3     # Sort documents by topic prevalence
4     top_docs = doc_topic_dist[:, topic_idx].argsort() [::-1] [:2]
5
6     print(f"\nTopic {topic_idx} - Top documents:")
7     for doc_idx in top_docs:
8         prevalence = doc_topic_dist[doc_idx, topic_idx]
9         print(f"    [{prevalence:.2%}] {documents[doc_idx]}")

```

Listing 3: Find most representative documents

#### Why Document-Based Interpretation?

Looking at actual documents is more reliable than top words because:

- Top words may be ambiguous or share meanings
- Documents provide context
- 95% prevalence means the document is almost entirely about that topic
- Reading the document tells you definitively what the topic represents

## 6 Non-negative Matrix Factorization (NMF)

#### Overview

NMF is a **deterministic** alternative to LDA. It uses linear algebra rather than probabilistic modeling to decompose documents into topics.

### 6.1 The Matrix Factorization Idea

#### NMF Decomposition

Given document-term matrix  $V$  (documents  $\times$  vocabulary):

$$V \approx W \cdot H$$

where:

- $W$ : Document-topic matrix (documents  $\times$  topics)
- $H$ : Topic-word matrix (topics  $\times$  vocabulary)
- All entries in  $W$  and  $H$  are  $\geq 0$  (non-negative)

## 6.2 Matrix Multiplication Review

### Matrix Multiplication Example

$$W = \begin{pmatrix} 1 & 2 \\ 3 & 4 \end{pmatrix}, \quad H = \begin{pmatrix} 1 & 1 & 0 \\ 0 & 0 & 2 \end{pmatrix}$$

$$V = W \cdot H = \begin{pmatrix} 1 \cdot 1 + 2 \cdot 0 & 1 \cdot 1 + 2 \cdot 0 & 1 \cdot 0 + 2 \cdot 2 \\ 3 \cdot 1 + 4 \cdot 0 & 3 \cdot 1 + 4 \cdot 0 & 3 \cdot 0 + 4 \cdot 2 \end{pmatrix} = \begin{pmatrix} 1 & 1 & 4 \\ 3 & 3 & 8 \end{pmatrix}$$

Rules:

- $(m \times n) \cdot (n \times p) = (m \times p)$
- Element  $(i, j)$  = dot product of row  $i$  from first matrix and column  $j$  from second

## 6.3 NMF for Topic Modeling

### Concrete NMF Example

Three documents with vocabulary [cats, dogs, bark, purr, growl]:

- Doc 1: “cats meow”  $\rightarrow [1, 0, 0, 1, 0]$
- Doc 2: “dogs bark”  $\rightarrow [0, 1, 1, 0, 0]$
- Doc 3: “cats purr dogs growl”  $\rightarrow [1, 1, 0, 1, 1]$

$$V = \begin{pmatrix} 1 & 0 & 0 & 1 & 0 \\ 0 & 1 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 1 \end{pmatrix}$$

NMF finds:

$$W = \begin{pmatrix} 1 & 0 \\ 0 & 1 \\ 0.5 & 0.5 \end{pmatrix}, \quad H = \begin{pmatrix} 1 & 0 & 0 & 1 & 0 \\ 0 & 1 & 1 & 0 & 1 \end{pmatrix}$$

Interpretation:

- Topic 1: cats-related (cats, purr)
- Topic 2: dogs-related (dogs, bark, growl)
- Doc 3: 50% Topic 1, 50% Topic 2

## 6.4 NMF vs LDA

Aspect	LDA	NMF
Approach	Probabilistic	Deterministic
Algorithm	EM (iterative)	Matrix factorization
Reproducibility	Stochastic (varies)	Deterministic (same result)
Interpretability	Often better	Harder to interpret $W$ values
Speed	Slower	Faster
Constraints	Probabilities sum to 1	Only non-negativity

### NMF Interpretation Challenge

NMF's  $W$  matrix entries are just non-negative numbers, not probabilities:

- Values like 2.0 and 5.0 are valid
- Harder to say “60% Topic 1”
- Requires normalization for probability interpretation

## 6.5 NMF Implementation

```

1 from sklearn.decomposition import NMF
2 from sklearn.feature_extraction.text import CountVectorizer
3
4 # Same document-term matrix X from before
5 nmf = NMF(n_components=2, random_state=42)
6 W = nmf.fit_transform(X) # Document-topic matrix
7 H = nmf.components_ # Topic-word matrix
8
9 # Display top words per topic
10 for idx, topic in enumerate(H):
11     top_words = [feature_names[i] for i in topic.argsort()[-5:]]
12     print(f"Topic {idx}: {top_words}")

```

Listing 4: NMF in Python

## 7 Choosing the Number of Topics

### Overview

Selecting the optimal number of topics is crucial. Two key metrics help: **coherence** (within-topic word relatedness) and **exclusivity** (between-topic word distinctiveness).

### 7.1 Coherence

#### Topic Coherence

Coherence measures how related the top words within a topic are to each other. Higher coherence means the topic's top words frequently co-occur in documents.

**Intuition:** If Topic 1's top words are [dogs, walks, fetch, leash, park], these words should appear together in documents more often than random word pairs.

#### Computing Coherence (Simplified)

For top words in a topic, compute co-occurrence:

$$\text{Coherence} \propto \sum_{i < j} \log \frac{P(w_i, w_j) + \epsilon}{P(w_i) \cdot P(w_j)}$$

This is essentially a log-transformed correlation measure.

## 7.2 Exclusivity

### Topic Exclusivity

Exclusivity measures how distinct topics are from each other. High exclusivity means top words in one topic don't appear as top words in other topics.

### Exclusivity Calculation (Simplified)

For word  $w$  in topic  $k$ :

$$\text{Exclusivity}(w, k) = \frac{P(w|\text{topic } k)}{\sum_{k'} P(w|\text{topic } k')}$$

If "dogs" has high probability only in Topic 1 and low in others, it has high exclusivity for Topic 1.

## 7.3 The Coherence-Exclusivity Trade-off

### Trade-off

Coherence and exclusivity often compete:

- More topics → Higher exclusivity but potentially lower coherence
- Fewer topics → Higher coherence but topics may overlap

### Important

#### Selection strategy:

1. Compute coherence and exclusivity for different numbers of topics
2. Standardize both metrics (z-scores)
3. Plot: x-axis = exclusivity, y-axis = coherence
4. Choose model closest to top-right corner (high both)

```

1 from gensim.models.coherencemodel import CoherenceModel
2
3 # After training LDA model
4 coherence_model = CoherenceModel(
5     model=lda_gensim,
6     texts=tokenized_docs,
7     dictionary=dictionary,
8     coherence='c_v',          # Coherence type
9     topn=20                  # Top 20 words per topic
10)
11 coherence_score = coherence_model.get_coherence()
12 print(f"Coherence: {coherence_score:.4f}")

```

Listing 5: Coherence calculation with gensim

## 7.4 Multiple Runs

### Stochastic Nature of LDA

LDA results vary between runs due to:

- Random initialization
- EM algorithm finding different local optima
- Topic labeling is arbitrary (Topic 1 in run A might be Topic 2 in run B)

**Best practice:** Run multiple times with different seeds, evaluate coherence/exclusivity, select best run.

## 8 LDA Implementation in R

### Overview

R has excellent support for topic modeling, especially for Structural Topic Modeling (STM). Learning basic R is worthwhile for NLP research.

### 8.1 Setting Up R

1. Download R from <https://cran.r-project.org/>
2. Download RStudio from <https://posit.co/>
3. Install R first, then RStudio (so RStudio finds R automatically)

### 8.2 Basic R Syntax for LDA

```

1 # Install and load packages
2 install.packages("topicmodels")
3 library(topicmodels)
4
5 # Create documents
6 documents <- c(
7   "cats are wonderful pets",
8   "cats and dogs are popular",
9   "dogs enjoy long walks",
10  # ... more documents
11 )
12
13 # Preprocessing
14 corpus <- Corpus(VectorSource(documents))
15 corpus <- tm_map(corpus, tolower)
16 corpus <- tm_map(corpus, removePunctuation)
17 corpus <- tm_map(corpus, removeWords, stopwords("english"))
18
19 # Create Document-Term Matrix
20 dtm <- DocumentTermMatrix(corpus)
21
22 # Fit LDA
23 lda_model <- LDA(dtm, k = 2, control = list(seed = 42))
24
25 # Get top terms per topic

```

```
26 terms(lda_model, 5)
```

Listing 6: LDA in R

## 8.3 R Markdown for Reports

### RMarkdown Files (.Rmd)

Similar to Jupyter notebooks:

- Mix code and text
- Code in “chunks” (like cells)
- **Ctrl+Enter** runs current line
- **Knit** creates HTML/PDF report
- Use # for sections, ## for subsections

## 9 Practical Applications

### 9.1 What Can You Do with Topic Models?

1. **Document Classification:** Assign documents to dominant topics
2. **Search:** Find documents about specific topics
3. **Trend Analysis:** Track topic prevalence over time
4. **Bias Detection:** Correlate topics with metadata (gender, source)
5. **Summarization:** Understand what a corpus is “about”

### 9.2 Real-World Example: Student Evaluations

#### Analyzing Student Evaluations

Study with 1 million student evaluations:

- Extracted 11 topics from evaluation text
- Topics: “caring instructor”, “interesting lectures”, “good feedback”, etc.
- Correlated with instructor gender

#### Findings:

- Female instructors: more mentions of “caring”, “facilitates discussion”, “nice feedback”
- Male instructors: more mentions of “humor”, “interesting”, “relevant”
- This pattern persisted after controlling for department and course type
- Suggests systematic bias in how students perceive instructors

### 9.3 Comparing Across Departments

#### Topic Variation by Division

Same student evaluation study:

- **Sciences:** “explains complex concepts effectively”
- **Humanities:** “facilitates effective discussions”
- **Freshman seminars:** “positive timely feedback”

These differences reflect genuine pedagogical differences across disciplines.

## 10 Structural Topic Modeling (STM) Preview

#### Overview

STM extends LDA by incorporating **covariates**—document-level metadata that can affect topic prevalence and content.

#### LDA vs STM

- **LDA:** Documents have topic proportions, but ignores metadata
- **STM:** Topic proportions can depend on covariates (author gender, publication source, date, etc.)

#### When STM Helps

Analyzing news articles with known sources:

- LDA: Discovers topics, then you manually correlate with sources
- STM: Directly models how source affects topic distribution
- STM produces more accurate topics by using all available information

#### Software Note

STM is primarily implemented in R (`stm` package). Python implementations exist but are unofficial. For serious STM work, learn R.

## 11 One-Page Summary

#### Summary

**Topic Modeling** discovers hidden themes in document collections.

#### Why Not Clustering?

- Documents contain multiple topics (soft assignment)
- Clustering forces hard assignment to one cluster

#### LDA (Latent Dirichlet Allocation):

- Probabilistic model: documents are mixtures of topics
- Topics are distributions over words
- Generative process: choose topic proportions, then for each word, choose topic then word
- Uses EM algorithm for parameter estimation
- Number of topics is a hyperparameter

#### NMF (Non-negative Matrix Factorization):

- Deterministic approach:  $V \approx W \cdot H$
- Faster but harder to interpret
- No probabilistic interpretation

#### Choosing Number of Topics:

- **Coherence:** Are top words related? (higher = better)
- **Exclusivity:** Are topics distinct? (higher = better)
- Balance both; maximize average of standardized scores

#### Interpreting Topics:

- Look at top words per topic
- Better: examine documents with highest topic prevalence
- Assign human-readable labels after analysis

#### Key Formulas:

$$\begin{aligned}\theta_m &\sim \text{Dirichlet}(\alpha) \quad (\text{topic proportions}) \\ z_n &\sim \text{Multinomial}(\theta_m) \quad (\text{topic selection}) \\ w_n &\sim \text{Multinomial}(\beta_{z_n}) \quad (\text{word selection}) \\ V &\approx W \cdot H \quad (\text{NMF decomposition})\end{aligned}$$

## 12 Glossary

### Key Terms

- **LDA:** Latent Dirichlet Allocation—probabilistic topic model
- **NMF:** Non-negative Matrix Factorization—deterministic topic model
- **STM:** Structural Topic Modeling—LDA with covariates
- **Dirichlet distribution:** Produces probability vectors summing to 1
- **Latent topic:** Hidden theme discovered from data (not predefined)
- **Topic proportions:** How much each topic contributes to a document

- **Coherence:** Metric measuring word co-occurrence within topics
- **Exclusivity:** Metric measuring distinctiveness between topics
- **EM algorithm:** Expectation-Maximization for latent variable models
- **MLE:** Maximum Likelihood Estimation—find parameters maximizing data probability
- **Undercomplete:** Bottleneck dimension < input dimension
- **Stacked autoencoder:** Deep autoencoder with multiple hidden layers
- **Bag of words:** Document representation ignoring word order
- **Document-term matrix:** Matrix of word counts (documents  $\times$  vocabulary)