

# CSCI E-89B Introduction to Natural Language Processing

Harvard Extension School

Dmitry Kurochkin

Fall 2025  
Lecture 5

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# Introduction to TF-IDF

- **Term Frequency-Inverse Document Frequency (TF-IDF)**

- ▶ TF-IDF is a statistical measure used to evaluate the importance of terms in documents relative to a corpus.

- ★ **Term Frequency (TF):** Measures how often a term appears in a document. In simplest form, TF is calculated as:

$$TF(\text{term}, \text{doc}) = \frac{\text{Number of times term appears in the document}}{\text{Total number of terms in the document}}$$

- ★ **Inverse Document Frequency (IDF):** Assesses how much information a word provides by downscaling terms that occur frequently across the corpus:

$$IDF(\text{term}) = \ln \left( \frac{\text{Total number of documents}}{\text{Number of documents containing the term}} \right)$$

- ▶ **TF-IDF:**

$$TF-IDF(\text{term}, \text{doc}) = TF(\text{term}, \text{doc}) \times IDF(\text{term})$$

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- **TF-IDF Example**
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# TF-IDF Example: Manual Computation

- **Documents:**

- ① Doc 1: "cat dog cat"
- ② Doc 2: "dog mouse dog"
- ③ Doc 3: "dog mouse"
- ④ Doc 4: "mouse cat dog"

- **Vocabulary:** {"cat", "dog", "mouse"}

- **Term Frequency (TF):**

- ▶ TF(cat): Doc 1 = 2/3, Doc 2 = 0/3, Doc 3 = 0/2, Doc 4 = 1/3
- ▶ TF(dog): Doc 1 = 1/3, Doc 2 = 2/3, Doc 3 = 1/2, Doc 4 = 1/3
- ▶ TF(mouse): Doc 1 = 0/3, Doc 2 = 1/3, Doc 3 = 1/2, Doc 4 = 1/3

- **Inverse Document Frequency (IDF):**

$$\text{IDF}(\text{cat}) = \ln(4/2) = 0.693$$

$$\text{IDF}(\text{dog}) = \ln(4/4) = 0$$

$$\text{IDF}(\text{mouse}) = \ln(4/3) = 0.288$$

# TF-IDF Example: Manual Computation (Continued)

- **TF-IDF Values:**

① Doc 1:

- ★ **cat**:  $2/3 \times \ln(4/2) = 2/3 \times 0.693 \approx 0.462$
- ★ **dog**:  $1/3 \times \ln(4/4) = 1/3 \times 0 = 0$
- ★ **mouse**:  $0/3 \times \ln(4/3) = 0/3 \times 0.288 = 0$

② Doc 2:

- ★ **cat**:  $0/3 \times \ln(4/2) = 0/3 \times 0.693 = 0$
- ★ **dog**:  $2/3 \times \ln(4/4) = 2/3 \times 0 = 0$
- ★ **mouse**:  $1/3 \times \ln(4/3) = 1/3 \times 0.288 \approx 0.096$

③ Doc 3:

- ★ **cat**:  $0/2 \times \ln(4/2) = 0/2 \times 0.693 = 0$
- ★ **dog**:  $1/2 \times \ln(4/4) = 1/2 \times 0 = 0$
- ★ **mouse**:  $1/2 \times \ln(4/3) = 1/2 \times 0.288 \approx 0.144$

④ Doc 4:

- ★ **cat**:  $1/3 \times \ln(4/2) = 1/3 \times 0.693 \approx 0.231$
- ★ **dog**:  $1/3 \times \ln(4/4) = 1/3 \times 0 = 0$
- ★ **mouse**:  $1/3 \times \ln(4/3) = 1/3 \times 0.288 \approx 0.096$

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# Advantages of TF-IDF over BoW

- Enhances Bag of Words (BoW) by incorporating global term significance.
  - ▶ **Limitations of BoW:** Simply counts occurrences, potentially leading common words to overshadow meaningful terms.
  - ▶ **Advantage of TF-IDF:** Adjusts term weights based on their occurrence in the corpus, offering a more comprehensive measure of term importance.
- Balances term frequency with significance across a corpus.
  - ▶ **Balancing Act:** Combines local document frequency with global significance, refining how insights are drawn.
  - ▶ **Applications:** Widely used for keyword extraction, document classification, and enhancing search engine results.
  - ▶ **Normalization:** Often employed to make TF-IDF scores comparable across documents.

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- **Modified Versions of TF-IDF**
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# Modified Versions of Term Frequency

- **Raw Term Count:**

- ▶  $f(t, d)$  represents the raw count of term  $t$  occurrences in document  $d$ .
- ▶ Can be affected by document length.

- **Term Frequency (TF):**

- ▶  $\text{TF}(t, d) = \frac{f(t, d)}{N_d}$
- ▶  $N_d$  is the total number of terms in document  $d$ .
- ▶ Default method in Scikit-learn's 'TfidfVectorizer'.

- **Logarithmically Scaled TF:**

- ▶  $\text{TF}_{\ln}(t, d) = 1 + \ln(f(t, d))$  if  $f(t, d) > 0$ ; else 0.
- ▶ Reduces the impact of high frequency terms.

- **Double Normalization (K):**

- ▶  $\text{TF}_{\text{norm}}(t, d) = K + (1 - K) \times \frac{f(t, d)}{\max\{f(w, d) : w \in d\}}$
- ▶ Balances term frequencies relative to the maximum term frequency in the document.
- ▶  $K$  is typically 0.5.

- **Boolean Term Frequency:**

- ▶  $\text{TF}_{\text{bool}}(t, d) = 1$  if the term appears in the document, otherwise 0.
- ▶ Ignores actual frequency, focusing on presence.

# Modified Versions of Inverse Document Frequency (IDF)

- **Standard IDF:**

- $\text{IDF}(t) = \ln\left(\frac{N}{\text{DF}(t)}\right)$ , where  $\text{DF}(t)$  is the document frequency of term  $t$ .
- High IDF for rare terms across documents.

- **Smoothed IDF:**

- $\text{IDF}_s(t) = \ln\left(\frac{1+N}{1+\text{DF}(t)}\right) + 1$
- Default method in Scikit-learn's 'TfidfVectorizer'.
- Includes smoothing to prevent division by zero.

- **Probabilistic IDF (ProbIDF):**

- $\text{IDF}_{\text{prob}}(t) = \ln\left(\frac{N-\text{DF}(t)}{\text{DF}(t)}\right)$
- Considers the probability of document exclusion.

- **Maximal IDF:**

- $\text{IDF}_{\text{max}}(t) = \ln\left(\frac{\max(\text{DF})}{1+\text{DF}(t)}\right)$
- Uses maximum document presence as a benchmark.

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# Hands-On: TF-IDF in Python Using Sklearn

```
from sklearn.feature_extraction.text import TfidfVectorizer

# Example text
documents = ["cat dog cat",
              "dog mouse dog",
              "cat mouse",
              "mouse cat dog"]

# Initialize TfidfVectorizer
vectorizer = TfidfVectorizer()

# Fit and Transform the documents
tfidf_matrix = vectorizer.fit_transform(documents)

# Display the Vocabulary and TF-IDF Representation
print("Vocabulary:\n", vectorizer.get_feature_names_out(), "\n")
print("TF-IDF Representation:\n", tfidf_matrix.toarray())
```

Vocabulary:

```
['cat', 'dog', 'mouse']
```

TF-IDF Representation:

```
[[0.81649658 0.40824829 0.          ]
 [0.          0.81649658 0.40824829]
 [0.70710678 0.          0.70710678]
 [0.40824829 0.40824829 0.40824829]]
```

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# Introduction to Word Embeddings

- **Concept:**

- ▶ Word embeddings are dense vector representations of words.
- ▶ They encode words into numerical vectors where semantically similar words have similar representation.
- ▶ Unlike traditional Bag of Words, embeddings capture context by considering the proximity of words in text.

- **Goal:**

- ▶ Aim to create a high-dimensional space where relationships and meanings between words reflect their context.
- ▶ Facilitate understanding of words in multiple dimensions, addressing ambiguity through context.
- ▶ Enable machines to understand and process text in ways similar to human cognition.

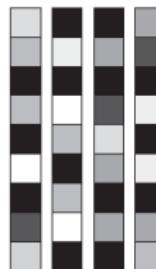
# Introduction to Word Embeddings

Word embeddings can be considered an alternative to one-hot encoding:



One-hot word vectors:

- Sparse
- High-dimensional
- Hardcoded

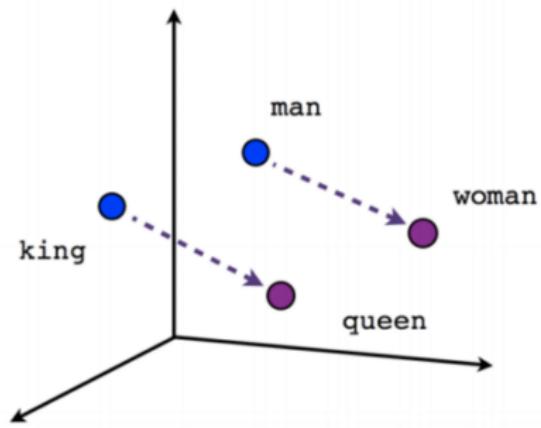


Word embeddings:

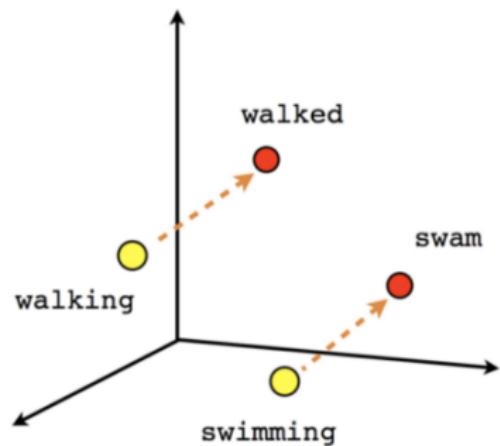
- Dense
- Lower-dimensional
- Learned from data

# Introduction to Word Embeddings

Example:



Male-Female



Verb tense

# Applications of Word Embeddings

- ① **Sentiment Analysis:** Improve accuracy by understanding context and nuances in opinions.
- ② **Machine Translation:** Capture word equivalencies across languages, enhancing translation quality.
- ③ **Information Retrieval and Search:** Enable intuitive and relevant search results by understanding related terms and concepts.
- ④ **Recommendation Systems:** Use vectors to determine item and user similarities, improving recommendations.
- ⑤ **Word Similarity and Analogy Tasks:** Enable finding similar or related words, and solve analogy tasks (e.g., "king" is to "queen" as "man" is to "woman").
- ⑥ **Text Classification:** Enhance accuracy by providing semantically rich word representations.
- ⑦ **Named Entity Recognition (NER):** Facilitate identifying and classifying named entities like people, organizations, and locations.

# Applications of Word Embeddings (Continued)

- ⑧ **Topic Modeling and Clustering:** Assist in grouping similar documents to uncover underlying themes within datasets.
- ⑨ **Question Answering Systems:** Improve retrieval and understanding of relevant information by capturing word relationships.
- ⑩ **Language Modeling and Generation:** Enhance sequence probability prediction, aiding in text understanding and generation.
- ⑪ **Relation Extraction:** Enable the extraction of semantic relationships between entities, enhancing databases and knowledge graphs.
- ⑫ **Social Media Monitoring:** Detect sentiment trends and conduct topic analysis in social media text.
- ⑬ **Spelling Correction and Typing Suggestions:** Provide more accurate recommendations by understanding context.

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings**
- Word2Vec
- GloVe

## 3 Practical Application in Python

# Limitations of Word Embeddings

- **Out-of-Vocabulary (OOV) Words:**

- ▶ Static embeddings cannot handle words not seen during training.
- ▶ Subword approaches like FastText can mitigate this.

- **Context Insensitivity (Static Embeddings):**

- ▶ Single vector representation per word fails to capture different meanings in diverse contexts.
- ▶ Modern context-specific embeddings address this limitation.

- **Bias Reflection and Amplification:**

- ▶ Embeddings can encapsulate biases present in training data.
- ▶ These biases can influence downstream applications, requiring debiasing techniques.

- **Resource Intensive:**

- ▶ Training and fine-tuning embeddings require significant computational resources.
- ▶ Infeasible for real-time deployment in limited-resource environments.

# Limitations of Word Embeddings (Continued)

- **Semantic Drift:**

- ▶ Static embeddings do not adapt to changes in language over time.
- ▶ Requires periodic retraining to maintain relevance with language evolution.

- **Low-Resource Language Limitations:**

- ▶ Less effective for languages with limited training data, potentially yielding poorer performance.

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- **Word2Vec**
- **GloVe**

## 3 Practical Application in Python

# Word2Vec: Introduction and Key Models

- **Developed by Google:**

- ▶ Created by a team at Google led by Tomas Mikolov in 2013.
- ▶ Aimed to efficiently process vast amounts of text to produce high-quality word embeddings.
- ▶ Revolutionized NLP by significantly improving the computational efficiency of training word vectors.

- **Key Models:**

- ▶ **Skip-gram:**

- ▶ Predicts surrounding context words for a given target word.
- ▶ Works well with small datasets.
- ▶ Effective for modeling infrequent words by leveraging context.

- ▶ **CBOW (Continuous Bag of Words):**

- ▶ Predicts a target word based on a given context of surrounding words.
- ▶ More computationally efficient than Skip-gram for large datasets.
- ▶ Tends to perform better with frequent words.

# Word2Vec: CBOW and Applications

- **Properties:**

- ▶ Utilizes shallow neural networks with one hidden layer.
- ▶ Trained using techniques like negative sampling or hierarchical softmax to optimize learning.
- ▶ Capable of capturing semantic relationships, enabling vector arithmetic (e.g., "king" - "man" + "woman" = "queen").

- **Applications:**

- ▶ Enhances document clustering, sentiment analysis, and recommendation systems.
- ▶ Provides foundational embeddings for more complex models like BERT and GPT.

# Contents

## 1 Term Frequency-Inverse Document Frequency (TF-IDF)

- Introduction to TF-IDF
- TF-IDF Example
- Advantages of TF-IDF over BoW
- Modified Versions of TF-IDF
- Hands-On: TF-IDF in Python Using Sklearn

## 2 Word Embeddings

- Introduction to Word Embeddings
- Limitations of Word Embeddings
- Word2Vec
- GloVe

## 3 Practical Application in Python

# GloVe: Introduction and Core Idea

- **Developed by Stanford:**

- ▶ Created by researchers at the Stanford NLP Group, including Jeffrey Pennington, Richard Socher, and Christopher D. Manning.
- ▶ Released in 2014 to address limitations in word embeddings by integrating statistical and contextual information.

- **Core Idea:**

- ▶ Utilizes a word co-occurrence matrix, recording how frequently words appear together across a corpus.
- ▶ This matrix serves as the foundation for learning embeddings that reflect both direct and indirect word relations.

- **Training Objective:**

- ▶ Aims to reconstruct the log-probabilities of word co-occurrences, where  $X_{ij}$  is the frequency at which word  $j$  appears in the context of word  $i$ .
- ▶ Optimizes embeddings so that the dot product between two word vectors approximates the logarithm of their probability of co-occurrence.
- ▶ Captures global context from the entire corpus while also focusing on local word relationships.

# GloVe: Training Objective and Advantages (Continued)

- **Advantages:**

- ▶ **Captures Complex Linguistic Patterns:**

- ★ GloVe embeddings encapsulate both semantic (meaning-related) and syntactic (grammar-related) relationships.
    - ★ Effective at distinguishing between words with similar meanings and different contexts through subtle variations captured by co-occurrence.

- ▶ **Holistic Text Understanding:**

- ★ Utilizes global statistical information by leveraging co-occurrence frequencies across the entire corpus.
    - ★ Facilitates understanding of broader language structure, supporting tasks like sentiment analysis with nuanced context comprehension.

- ▶ **Semantic Arithmetic with Embeddings:**

- ★ Supports intuitive operations, such as vector arithmetic:  
"king" - "man" + "woman" = "queen".
    - ★ Demonstrates the ability to perform analogy tasks, reflecting the interconnected semantic space mapped by embeddings.

# Comparison: Word2Vec vs. GloVe

- **Word2Vec:**
  - ▶ Uses local context.
  - ▶ Faster training, lower memory footprint.
- **GloVe:**
  - ▶ Combines global statistical information with local context.
  - ▶ More effective for capturing deeper patterns.
- **Selecting Models:**
  - ▶ Choice depends on dataset size and task specifics.
  - ▶ Word2Vec for large-scale, real-time adjustments.
  - ▶ GloVe for in-depth semantic understanding.

# Practical Application in Python

```
from gensim.models import Word2Vec
from glove import Glove, Corpus

# Sample Data
sentences = [["cat", "sat", "on", "the", "mat"],
              ["dog", "barked"],
              ["cat", "chased", "dog"]]

# Word2Vec Model
print("Training Word2Vec model...")
word2vec_model = Word2Vec(sentences, vector_size=100, window=5, min_count=1, workers=4)
print("Word2Vec model trained.")

# Accessing Word2Vec vector
print("Word2Vec Vector for 'cat':")
print(word2vec_model.wv['cat'])

# Preparing corpus for GloVe
corpus = Corpus()
print("Fitting GloVe corpus...")
corpus.fit(sentences, window=10)

# GloVe Model
print("Training GloVe model...")
glove_model = Glove(no_components=100, learning_rate=0.05)
glove_model.fit(corpus.matrix, epochs=30, no_threads=4, verbose=True)
print("GloVe model trained.")

# Accessing GloVe vector
glove_model.add_dictionary(corpus.dictionary) # Aligning GloVe with corpus
print("GloVe Vector for 'cat':")
print(glove_model.word_vectors[corpus.dictionary['cat']])
```

# Practical Application in Python

```
from gensim.models import Word2Vec
from glove import Glove, Corpus

# Sample Data
sentences = [["cat", "sat", "on", "the", "mat"],
              ["dog", "barked"],
              ["cat", "chased", "dog"]]

# Word2Vec Model
word2vec_model = Word2Vec(sentences, vector_size=100, window=5, min_count=1, workers=4)

# Preparing corpus for GloVe
corpus = Corpus()
corpus.fit(sentences, window=10)
glove_model = Glove(no_components=100, learning_rate=0.05)
glove_model.fit(corpus.matrix, epochs=30, no_threads=4, verbose=True)
```