

CSCI E-103

*Data Engineering for Analytics to Solve Business Challenges*

# Towards Reproducible Machine Learning

*Lecture 08*

Anindita Mahapatra & Eric Gieseke

Harvard Extension, Fall 2025

# Agenda

- ML Review for the Data Persona
- Role of Data Engineering in ML
- What is a Model Lifecycle
- Challenges in ML development
- ML Pipelines - Formalizing the ML development process for Model Reproducibility
- Feature Store
- MLOps - Managing ML lifecycle using SDLC discipline
- Intro to MLFlow
- Lab
  - MLFlow for model lifecycle management & MLOps
  - <https://docs.databricks.com/aws/en/notebooks/source/mlflow/mlflow-classic-ml-e2e-mlflow-3.html>
    - Replace main.default with cscie103\_catalog.default

# Review Previous Material

Scale Out refers to

Autoscale uses

Ability to use the service anytime is

Planning to switch to a different region on an outage

Ability to grow/shrink nodes refers to

Example of a Bridge Pattern in Spark

What are the advantages of multi-hop pattern

What is BI Vs BA

2 KPIs of BI tool performance

Architecture paradigm that supports AI+BI

Adding more nodes

Scale out

Availability

Disaster Recovery

Elasticity

Connector Pattern

Cater to different SLAs, more robust, modular

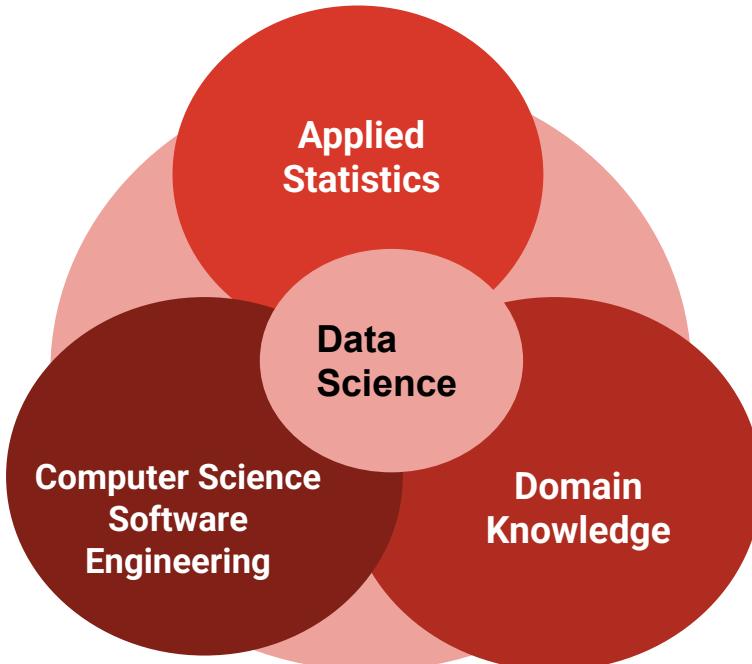
What Happened Vs Why/What Next

Latency & Concurrency

Lakehouse

# Data Science

## Process & Skills



Using Data to Solve  
Measurable Real World Problems

“The scientific method is a **dynamic** process that allows for the **continual advancement** of human understanding of the natural world.

It is characterized by its commitment to empirical evidence, objectivity, and the use of critical thinking to refine and expand our knowledge.”

# Algorithms & Machine Learning

## How it relates to Data Science

### Rule Based

Set of facts + Set of Rules

Deterministic - if - then - else

Beyond a certain set, it is unmanageable

### AI Based

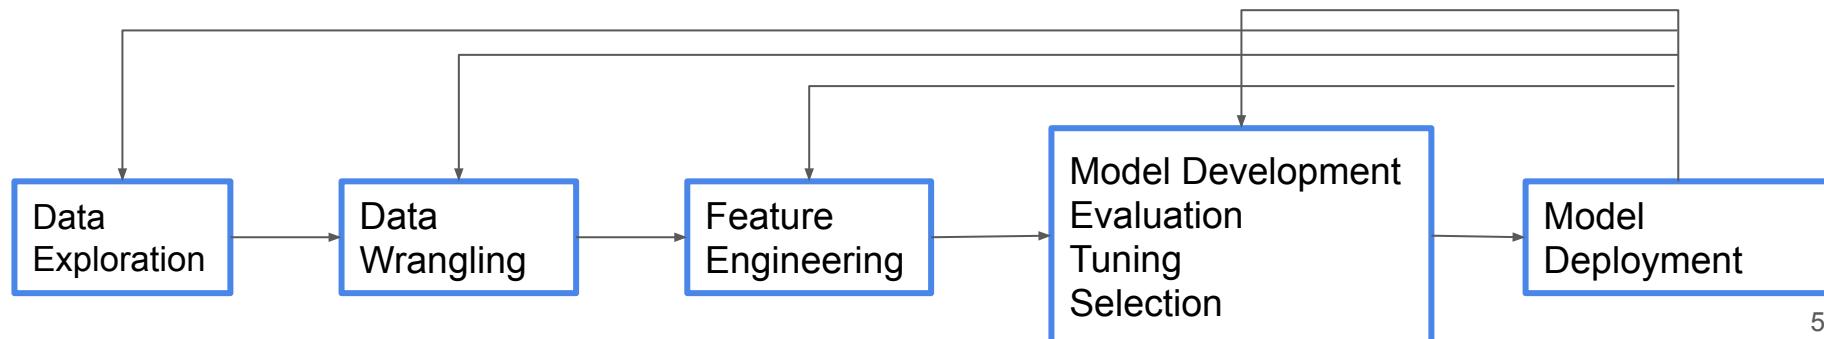
May have a set of facts

Probabilistic

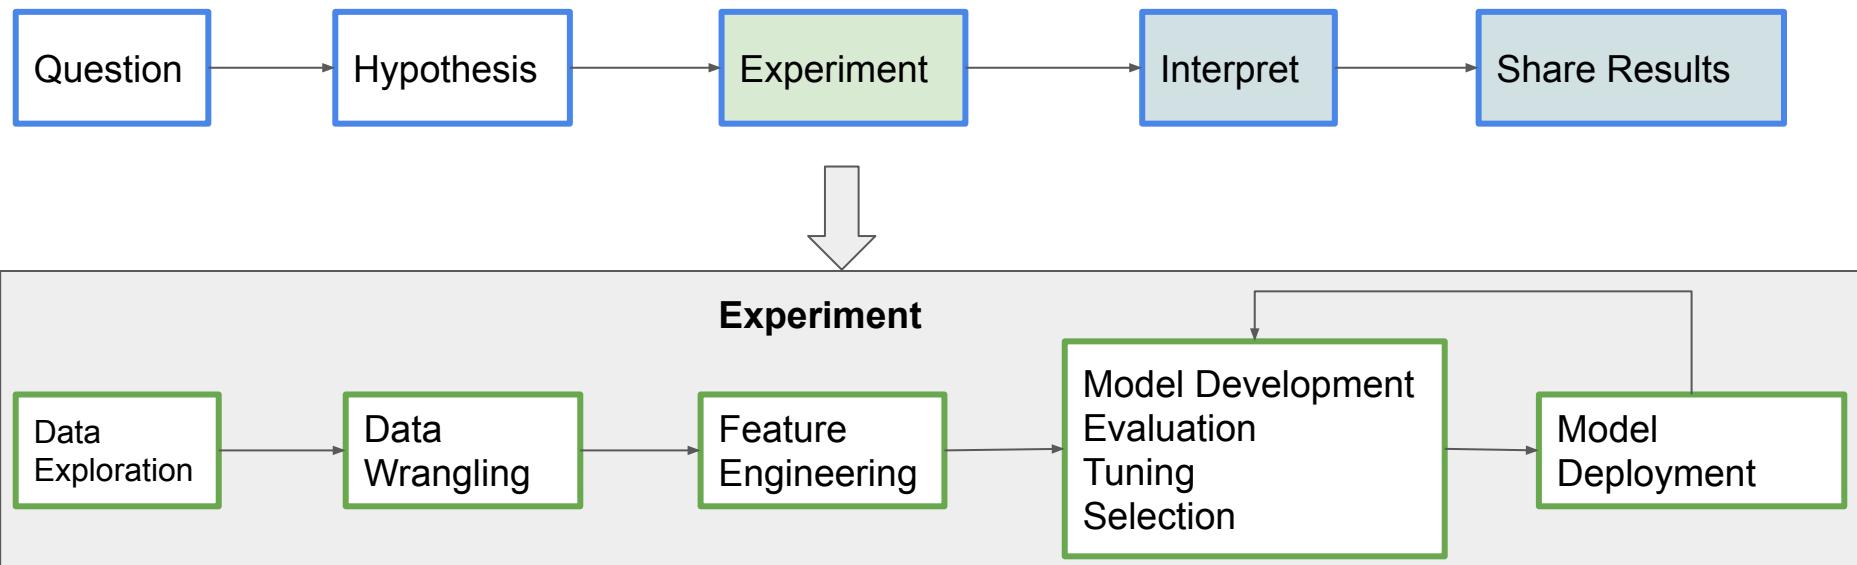
Require more data points and are scalable

**AI/ML:** Perform tasks without being explicitly told via rules aka if/then/else i.e. learn relationships (algorithms)

- Columns -> Features -> Predict
- More data -> more wisdom -> better
- More data -> algorithm needs to be computationally efficient



# ML as a part of Data Science



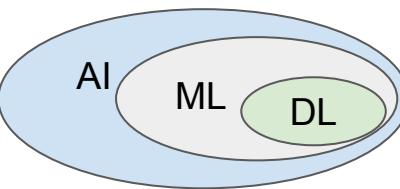
ML is a part of Data Science

Primarily the Experiment and Analysis Phase

Extends into Interpretability and Result Delivery

# AI Vs ML Vs DL

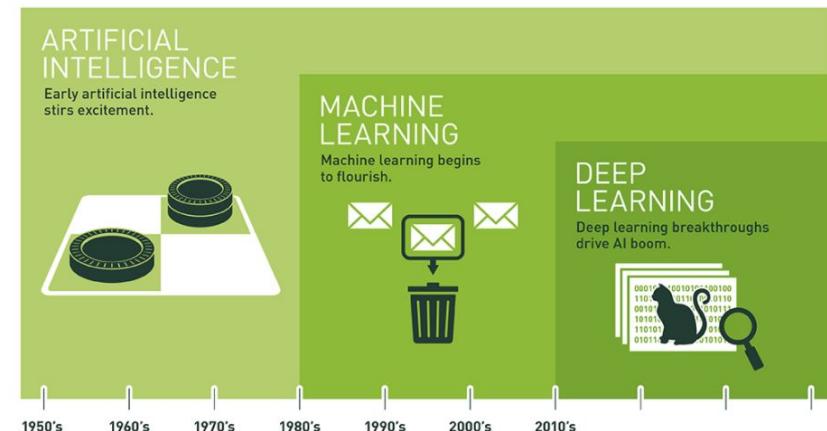
## Role of Spark



**AI** - machines perform tasks normally requiring human intelligence

**ML** - the process of learning from data without being explicitly programmed aka rule based

**DL** - Deep Learning: complex neural networks detect patterns in large unstructured data sets



*"Hidden Technical Debt in Machine Learning Systems", Google*

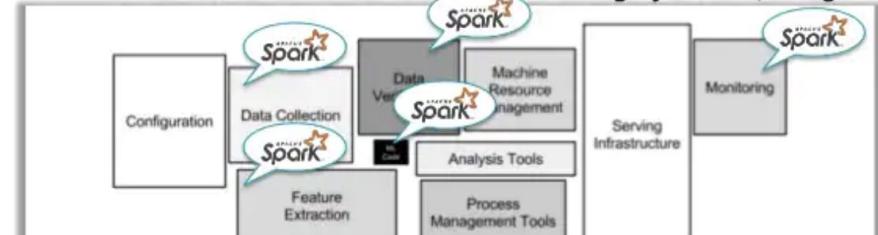


Figure 1: Only a small fraction of real-world ML systems is composed of the ML code, as shown by the small black box in the middle. The required surrounding infrastructure is vast and complex.

FLEXIBLE

FAST

BIG DATA

So what has changed lately?

- Lots of Data & Affordable Compute

# ML Frameworks & Libraries

- Classical/Statistical Algorithms
  - Suitable for most DS problems around structured/semi data
  - Mature and work with low data volume
- Neural Network/Deep Learning
  - Mostly unstructured data
  - Typically, performance increases significantly with training data volume
- Generative AI
  - LLM Large Language Models (text)
  - GAN Generative Adversarial Networks (images)

AutoML  
Fast & Easy but less customizable

Notebook based  
Powerful, but more complex

# Classification examples of ML types

	<i>Supervised Learning</i>	<i>Unsupervised Learning</i>
<i>Discrete</i>	classification or categorization	clustering
<i>Continuous</i>	regression	dimensionality reduction

How much will it rain tomorrow?

Will it rain tomorrow?

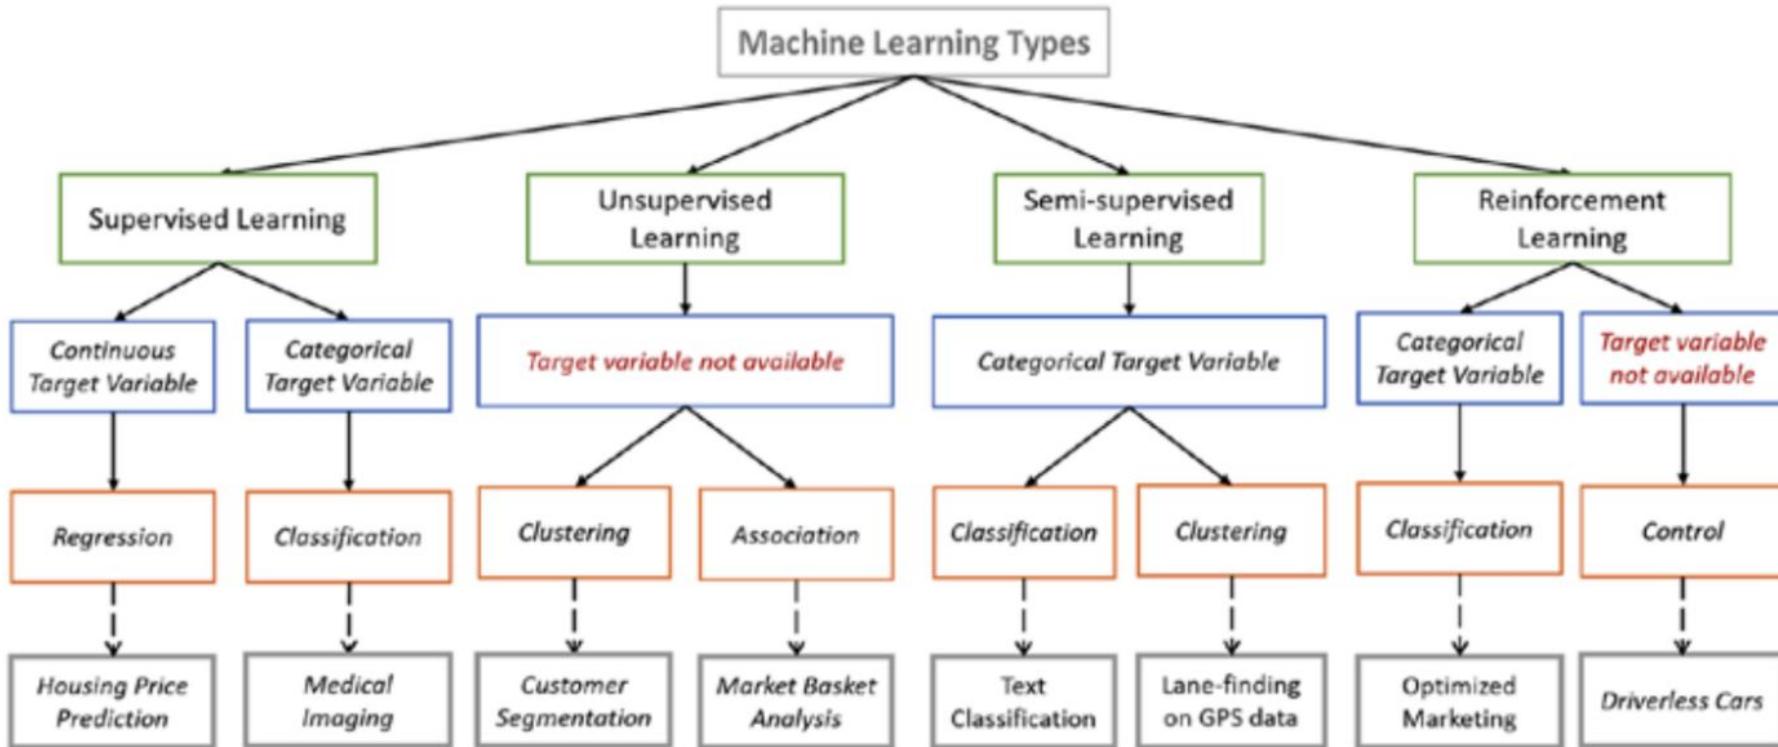
Which days of the week can be grouped based on weather patterns?

What are the key factors that influence consumer purchasing?

# Art of Prediction

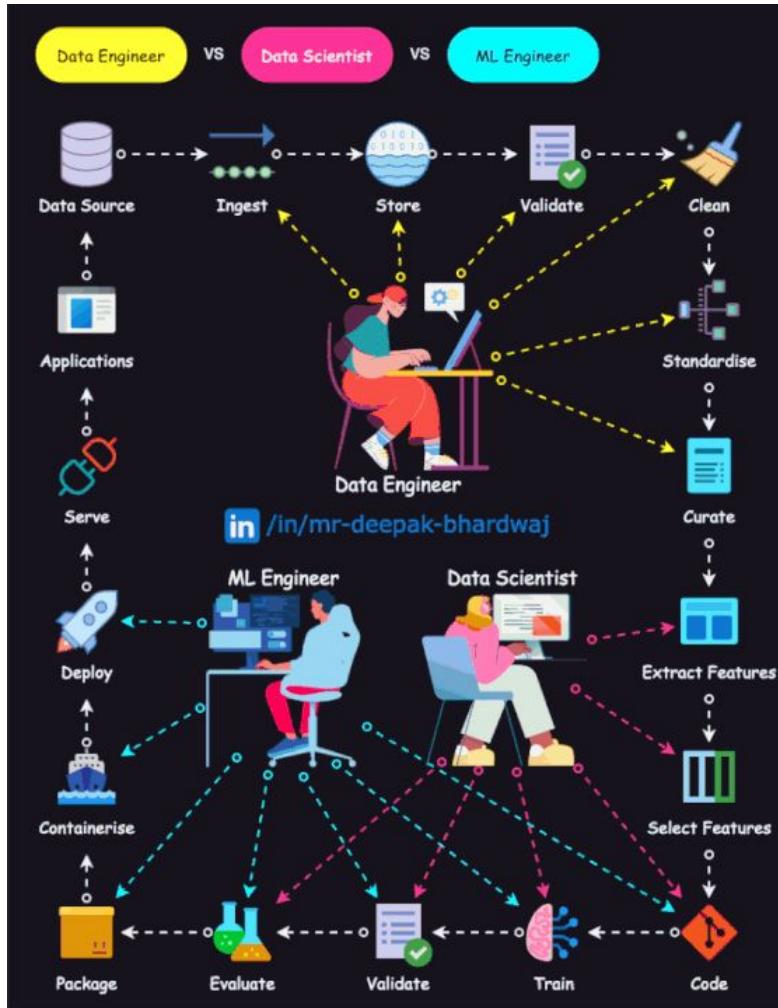
The art of looking into the future using the data from the past with/without labels

Lots of data & simple algorithms are more effective than the converse



# Role of Data Engineering in ML

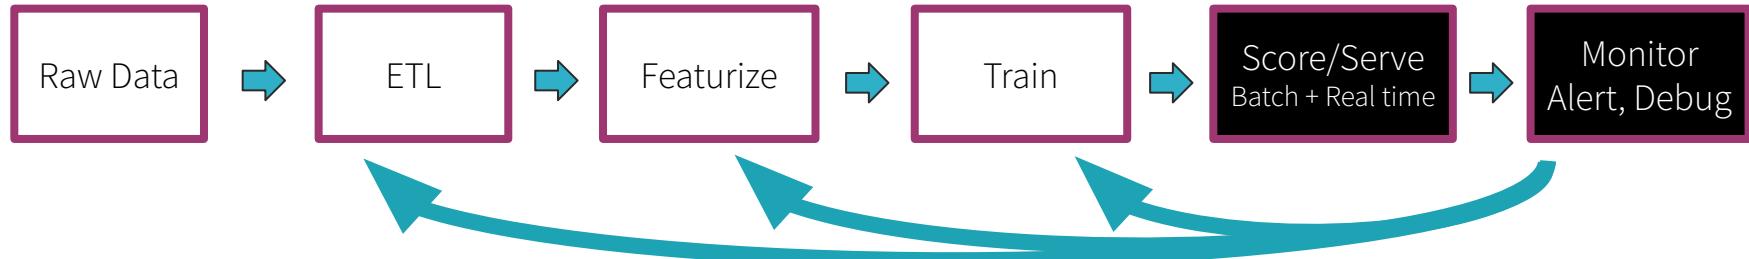
- Role Interplay
  - Data Engineer: Infrastructure Architects
  - Data Scientist: Insight Alchemists
  - ML Engineer: Automation Crafters
- Data Integration
  - Data engineering is responsible for merging disparate sources of data into a unified view that can be used by an AI model.
- High Quality Data
  - It ensures that high-quality data is available, accessible, and usable for AI models.
  - Feature Engineering
- Scalability
  - Data engineering ensures that the entire data infrastructure can scale to handle increasing amounts of data
- Monitoring
  - Data Drift detection signals model retraining



# Aspirations of a ML Platform

- Support & aid development on diverse ML tools & frameworks
  - Libraries
  - Languages
  - Collaboration
- Support various deployment strategies
  - Batch, Streaming, REST EndPoint, Edge
- Single node & distributed algorithms (CPU/GPU/Classical/Deep Learning)
- Support higher efficiencies
  - Faster training, tuning, tracking
- Feature Store for feature reuse
- AutoML for Citizen data scientists
- Monitor drift
- Aid in operationalizing ML
- Reproducible ML

# ML Lifecycle and Challenges



Tuning

Deploy

Model Mgmt

Collaboration

Scale

Governance

Feature Repository

Experiment Tracking

AutoML,  
Hyper-p. search

Remote Cloud Execution

Project Mgmt  
(scale teams)

Model Exchange

A/B Testing

CI/CD/Jenkins  
push to prod

Orchestration  
(Airflow, Jobs)

Lifecycle mgmt.

Data Drift

Model Drift

# From Data Ingestion to Model Deployment

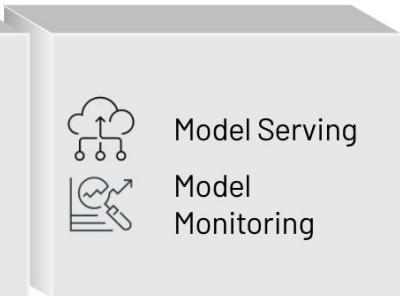
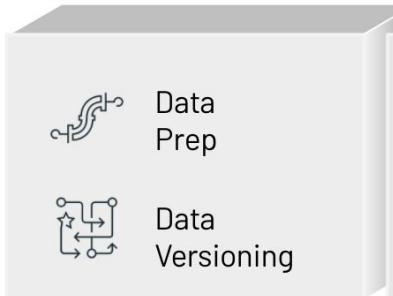
## Data prep designed for ML



## Out-of-the-box environment for all ML frameworks



## Deploy anywhere at any scale



Automation and Governance

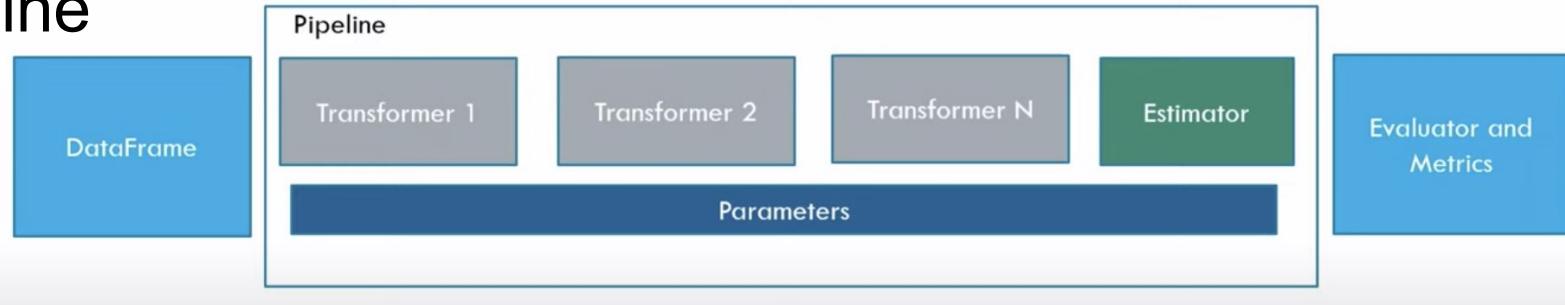
DataOps

ModelOps

DevOps

# ML Pipeline

## Stages



- **Transformers** (preprocessing)
  - String Indexer : string to index
  - One Hot Encoding : index to a binary vector
  - Standard Scaler : Scales the data to unit standard deviation.
  - Imputer: completes missing values in a dataset, either using the mean or the median
  - VectorAssembler: combines a given list of columns into a single vector column
- **Estimators** (learn)
  - Fit: Input is a dataframe, output is a model

**Metrics:** Evaluation metrics explain the performance of a model

**Evaluators:** functions to calculate performance metrics based on actuals & predicted values

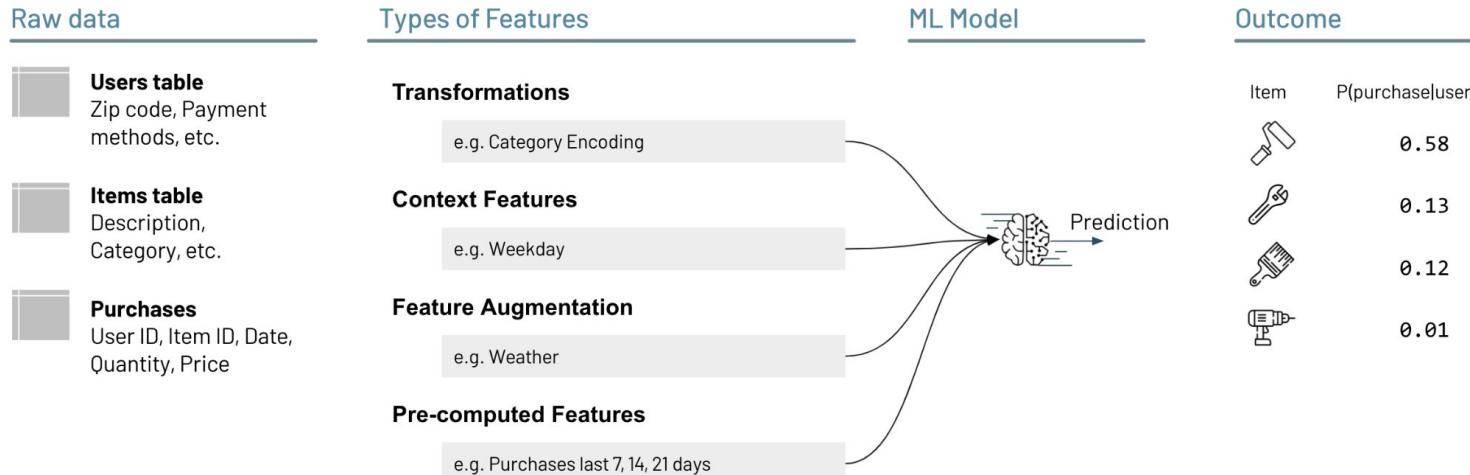
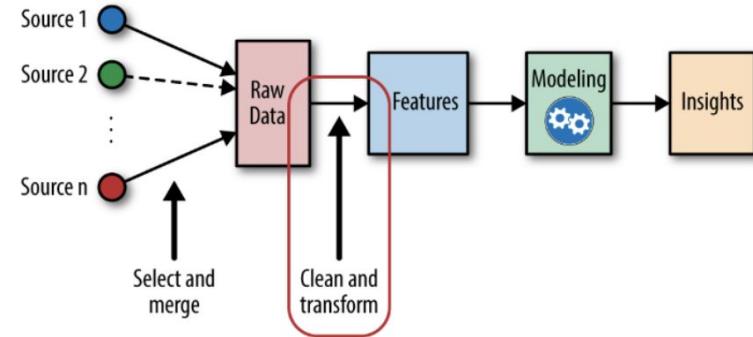
## Confusion Matrix

	Actually Positive (1)	Actually Negative (0)
Predicted Positive (1)	True Positives (TPs)	False Positives (FPs)
Predicted Negative (0)	False Negatives (FNs)	True Negatives (TNs)

# Feature Engineering

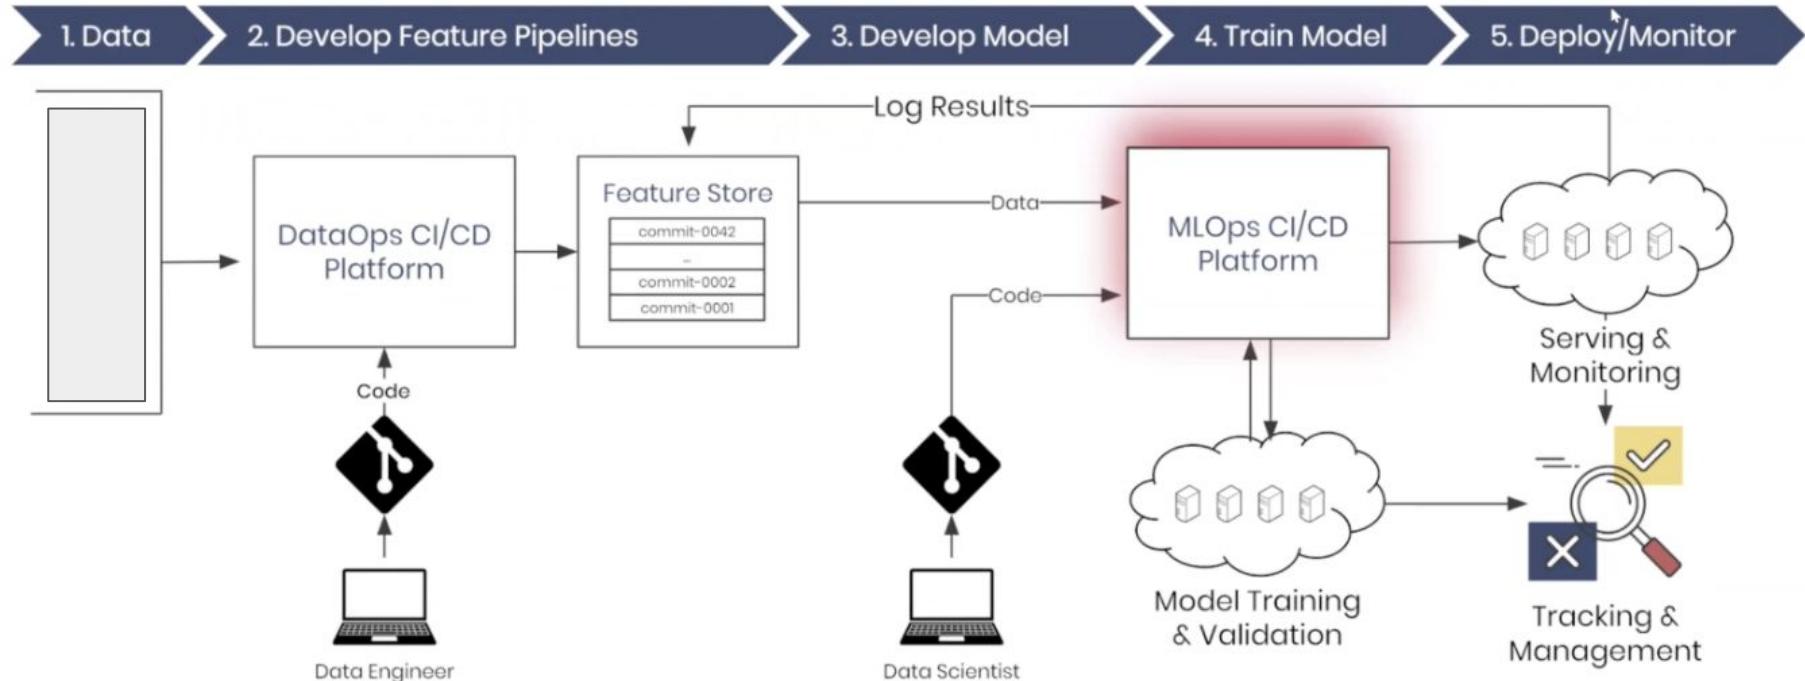
## Feature Store

Process of using domain knowledge of the data to create features that make machine learning algorithms work.



*Coming up with features is difficult, time-consuming, requires expert knowledge.  
“Applied machine learning” is basically feature engineering.*  
— [Andrew Ng](#), Machine Learning and AI via Brain simulations

# MLOps



# MLflow

Open source project

Conventions, specs, tools

CLI, libraries, REST server, UI

Community backed

Framework and tool agnostic,  
on-prem & in-cloud

1) **API first**: built around REST APIs that allows submitting runs, models, etc. from any library & language

2) **Modular** : independent components (Tracking/Projects/Models) E.g., use MLflow's project format but not its deployment tools. Easy to integrate into existing ML platforms & workflows

3) **Easy** : minimal, ubiquitous dependencies. Runs the same way anywhere (local, cloud platforms) Easy for a single dev to use locally, or very large teams

4) **Ecosystem first** : choose open and flexible core abstractions and interface points (e.g. entry point to an MLProject is a shell script so no bias to language runtime) Make it easy to use ecosystem frameworks, rather than competing with those frameworks.

5) **Open Source & Open interface**: MLflow is designed to work with any ML library, algorithm, deployment tool or language.

## mlflow

### Tracking

Record and query experiments: code, metrics, parameters, artifacts, models

## mlflow

### Projects

Packaging format for reproducible runs on any compute platform

## mlflow

### Models

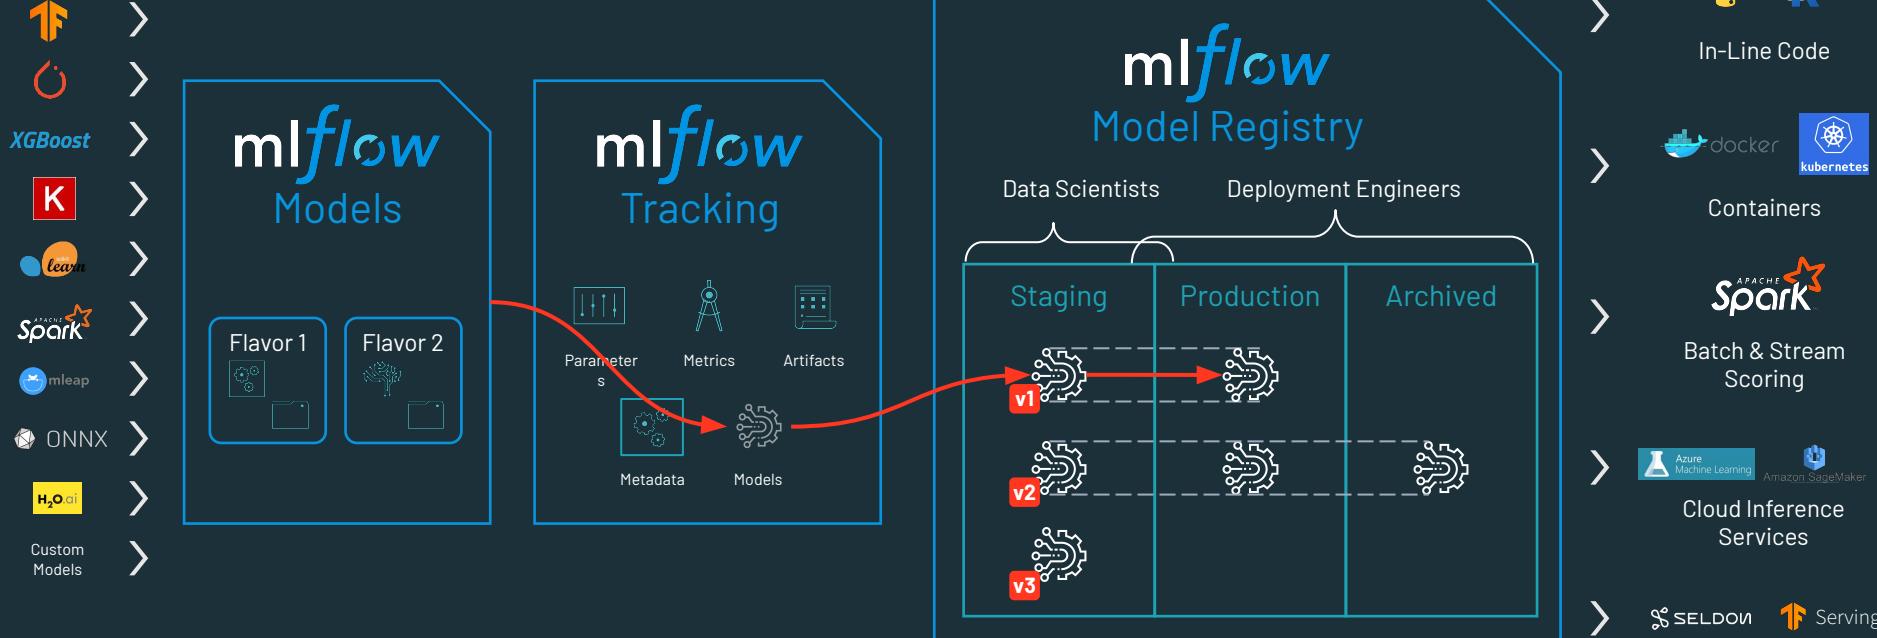
General model format that standardizes deployment options

## mlflow

### Model Registry

Centralized and collaborative model lifecycle management

# mlflow Model Lifecycle



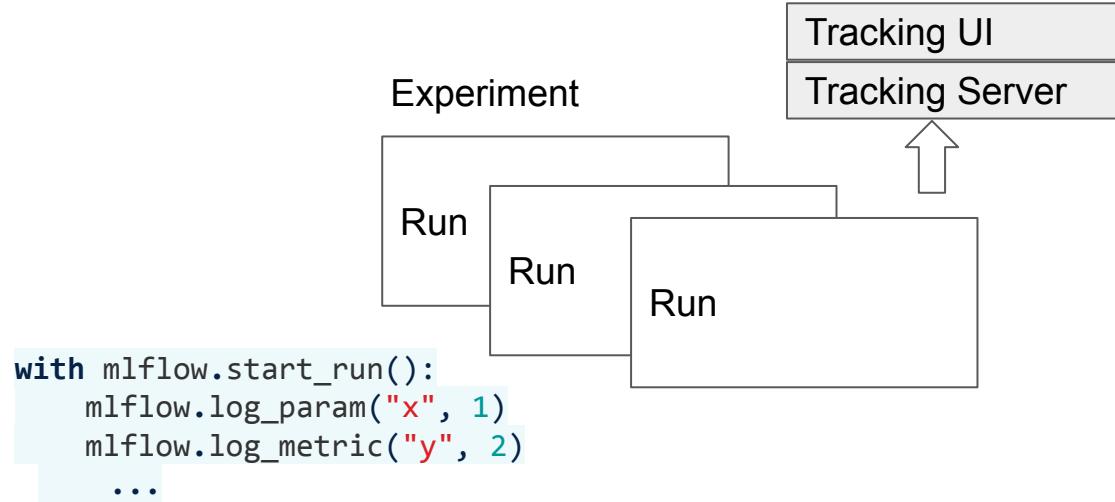
# MLFlow

## Tracking

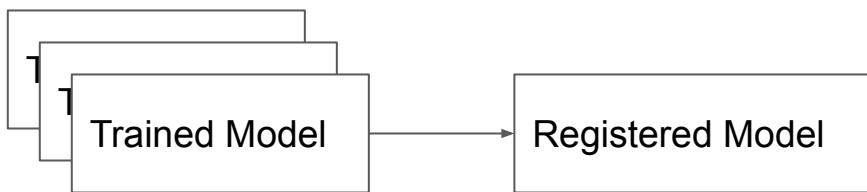
- `mlflow.start_run()`
- `mlflow.log_param()`
- `mlflow.log_metric()`
- `mlflow.set_tag()`
- `mlflow.log_artifact()`
- `mlflow.autolog()`

## Registry

- `create_registered_model()`
- `create_model_version()`
- `Transition_model_version_stage()`
- `list_registered_models()`



```
from mlflow.tracking import MlflowClient
client = MlflowClient()
client.create_registered_model("random-forest-model")
```



version      stage

# Model & Data Drift

What to monitor?

- Basic summary statistics of features & target
- Distribution of features & target
- Model Performance Metrics
- Business Metrics

Talks:

[Testing ML Models in Production](#)

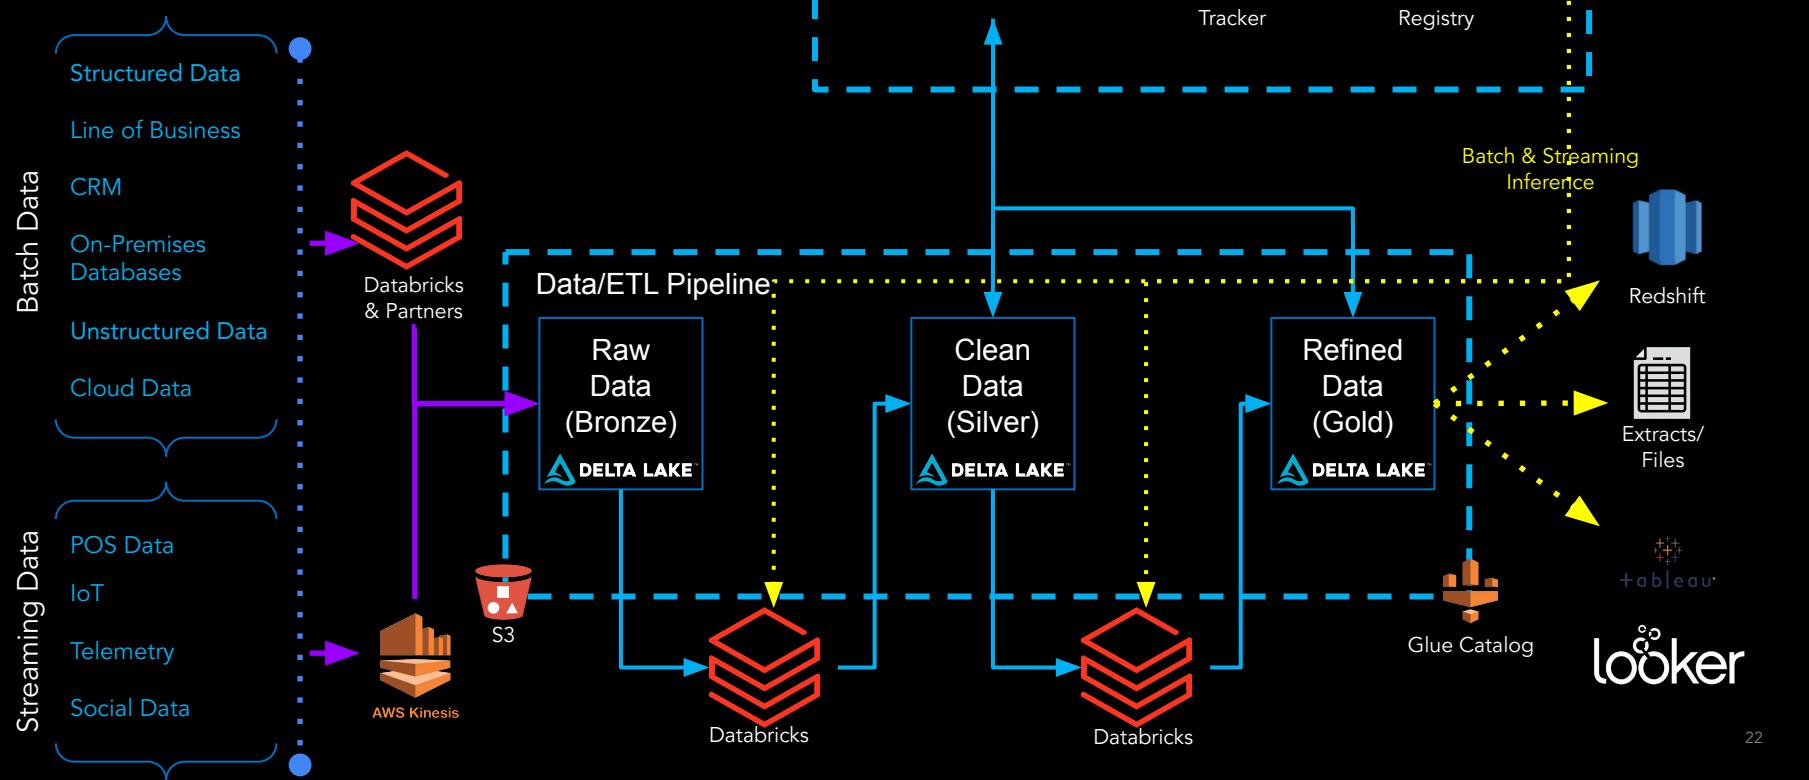
Webinars:

[Productionalizing ML](#)



Drift Type Identified	Action
Feature Drift	<ul style="list-style-type: none"><li>• Investigate feature generation process</li><li>• Retrain using new data</li></ul>
Label Drift	<ul style="list-style-type: none"><li>• Investigate label generation process</li><li>• Retrain using new data</li></ul>
Prediction Drift	<ul style="list-style-type: none"><li>• Investigate model training process</li><li>• Assess business impact of change in predictions</li></ul>
Concept Drift	<ul style="list-style-type: none"><li>• Investigate additional feature engineering</li><li>• Consider alternative approach/solution</li><li>• Retrain/tune using new data</li></ul>

# Where does MLflow fit in the architecture



# ToDos

- Read chapters 7,8, 9
- Lab
  - XGBoost: [Link](#)
  - Replace main.default with cscie103\_catalog.default or something that you have access to
- Lab
  - Scikit: [Link](#)
  - Update CATALOG\_NAME = "cscie103\_catalog"
  - Use local Hyperopt trials instead of SparkTrials.