

CSCI E-103

*Data Engineering for Analytics to Solve Business Challenges*

# Good Data -> Good Models

## Improving Model Insights

*Lecture 09*

Anindita Mahapatra & Eric Gieseke

Harvard Extension, Fall 2025

# Agenda

- Announcements
- Review Last Lecture
- MLOps
  - Model Creation, Selection, Registration, Deployment
  - From model centric to data centric
  - MLOps Maturity Progression
  - Automation as the tool to control model management
- Lab
  - MLOps
  - `%pip install dbdemos`
  - `import dbdemos`
  - `dbdemos.install('mlops-end2end', catalog='cscie103_catalog', schema='mlops', create_schema=True)`
  -

# Announcements

- Case Study #1 (Data Architecture Evolution Research)
- Assignment-4
  - Based on MLFlow
  - Released soon (we will review on Thursday)
- Case Study #2 (Industry Competitive Research)
  - Group
- Final Project
  - Group
- Quiz
  - Quiz#1 - done
  - Quiz#2 - released towards end of Nov

# Review Previous Material

Model staleness over time is referred to as

MLFlow is an open source framework for

MLFlow Tracking Server tracks

MLFlow Registry helps with model

Can MLflow work on-prem

Process of getting data ready for input to a model

Different types of model serving/scoring

Overfit refers to

What aids with feature reuse

AutoML refers to

Model Drift

Model lifecycle management

Parameters, metrics, artifact, code, data, tags, ..

Discoverability, versioning, promotion

Yes

Feature Engineering

Batch, streaming, realtime model serving

A statistical model that contains more parameters than can be justified by the data

Feature Store

Automating the tasks of applying machine learning to real-world problems

# People and process



**Business Stakeholder**

→ Responsible for business value of the ML solution



**Data Engineer**

→ Builds data pipelines



**Data Scientist**

→ Translates business problem; trains, tunes model



**ML Engineer**

→ Deploys ML model to production



**Data Governance Officer**

→ Responsible for data governance and compliance

# People and process



# Business Context: Customer Retention

No data science and ML can start without a business problem at hand.

You are on a marketing analytics team and you have a lot of **demographic** and **historical service data** on your customers that have **churned**, which has been put into a SQL Analytics dashboard.

The data team has been asked by business stakeholders if you can go further and **predict** which customers are likely to churn. Knowing this will allow the business to take action and **retain revenue**.

Sounds simple enough. What steps do we need to take?

[Example](#)

## Operational Workflow At a Glance:

Data Prep & Featurization

Baseline Model

Setup Webhooks

Promote Best Run to Registry

Testing

Staging Batch Inference

Update Dashboard

Schedule Monthly Retrain Job



Data Engineer



Data Scientist



ML Engineer

Operational Workflow:  
**Data Engineering**



**Data Engineer**

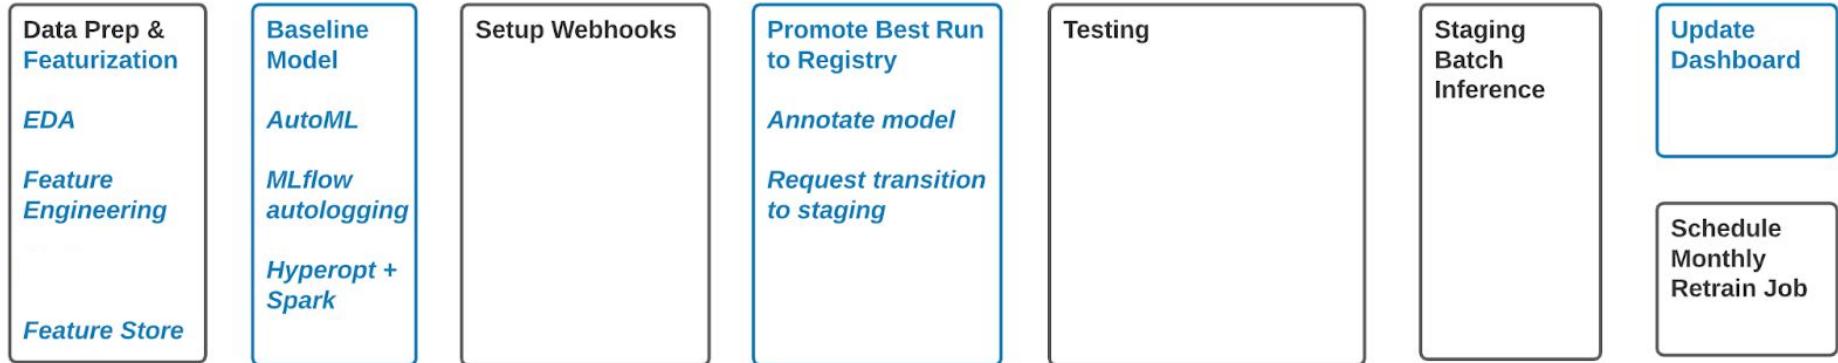


**Data Scientist**



**ML Engineer**

## Operational Workflow: Data Scientist



**EDA:** Exploratory Data Analysis



Data Engineer

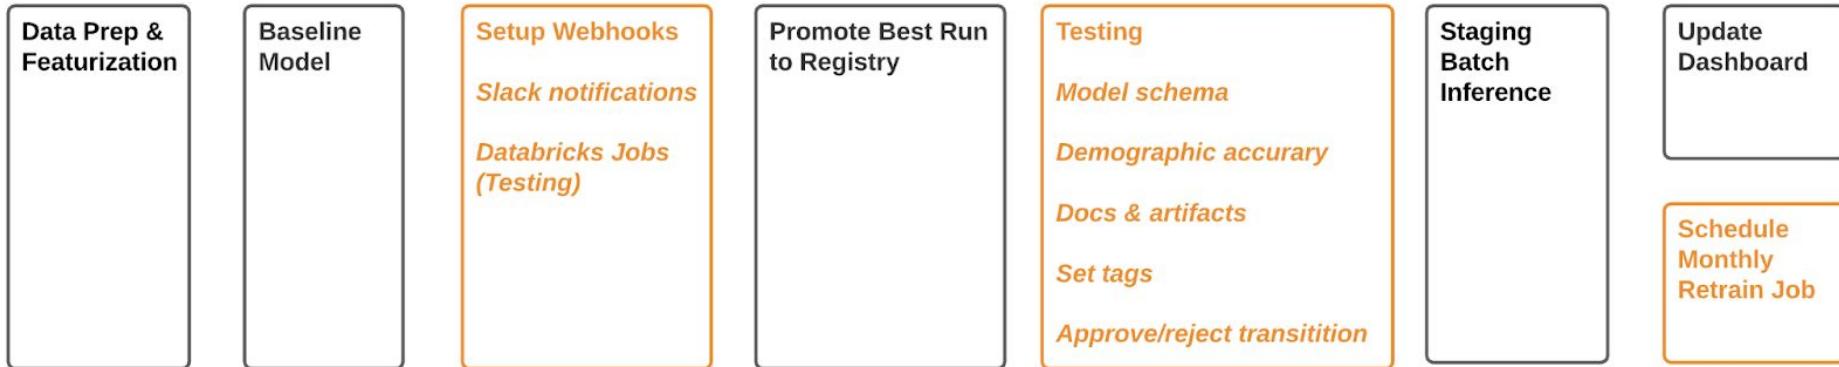


Data Scientist



ML Engineer

## Operational Workflow: ML Engineering



Data Engineer



Data Scientist



ML Engineer

## Operational Workflow: Data Team

**Data Prep & Featurization**

*ETL + EDA*  
*Feature Engineering*  
*Koalas*  
*Feature Store*

**Baseline Model**

*AutoML*  
*MLflow autologging*  
*Hyperopt + Spark*

**Setup Webhooks**

*Slack notifications*  
*Databricks Jobs (Testing)*

**Promote Best Run to Registry**

*Annotate model*  
*Request transition to staging*

**Testing**

*Model schema*  
*Demographic accuracy*  
*Docs & artifacts*  
*Set tags*  
*Approve/reject transition*

**Staging Batch Inference**

*Spark UDF*

**Update Dashboard**

*Schedule Monthly Retrain Job*



**Data Engineer**

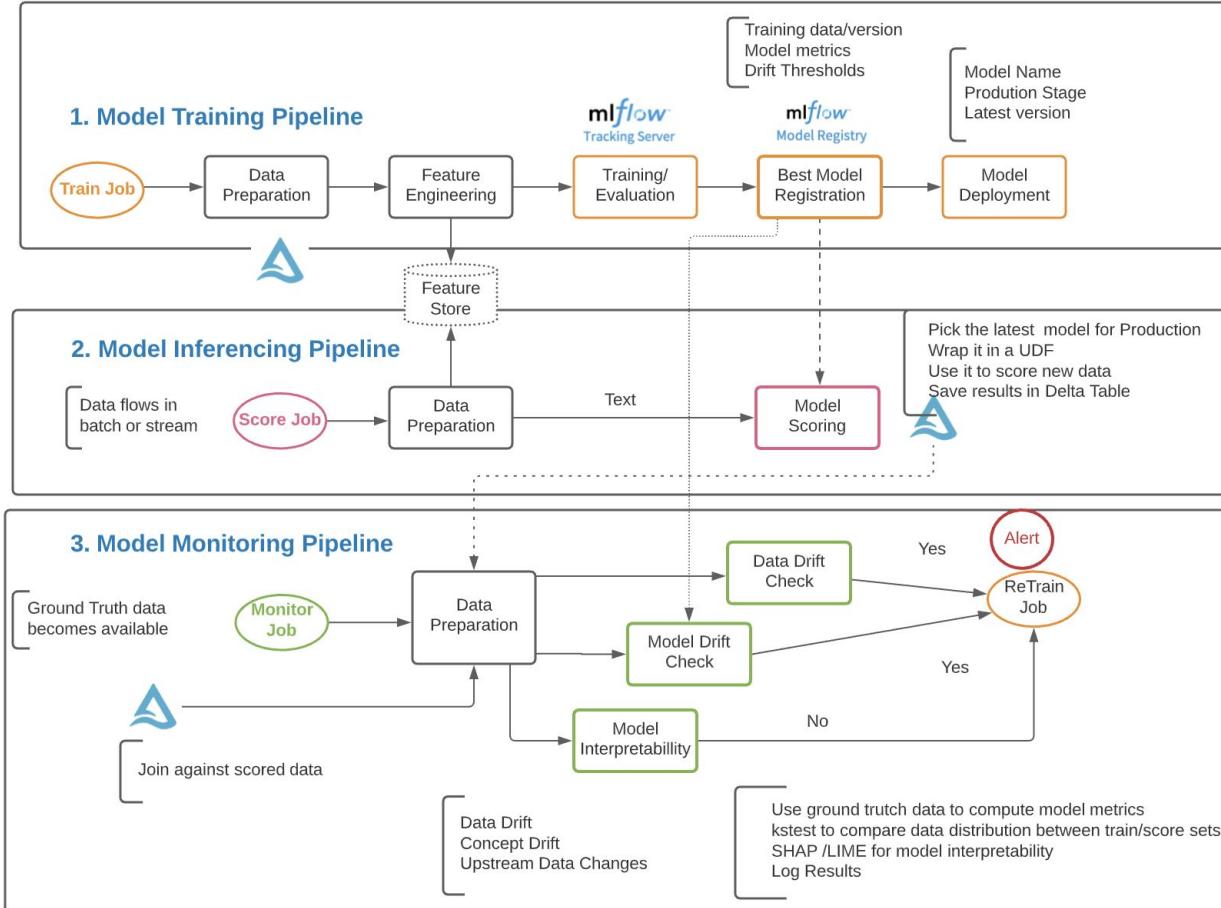


**Data Scientist**



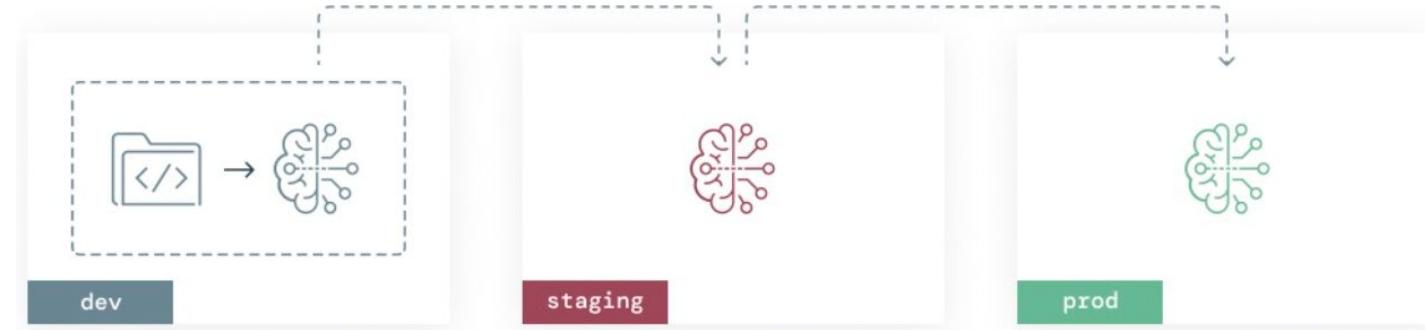
**ML Engineer**

# Model Training, Inferencing, and Monitoring



# ML deployment patterns

## Deploy Models



## Deploy Code



# Benefits of deploy code

<b>Automation</b>	↑ Supports automated retraining in locked-down env.
<b>Data access control</b>	↑ Only prod env needs read access to prod training data.
<b>Reproducible models</b>	↑ Eng control over training env, which helps to simplify reproducibility.
<b>Support for large projects</b>	↑ This pattern forces the DS team to use modular code and iterative testing, which helps with coordination and development in larger projects.
<b>Data science familiarity</b>	↓ DS team must learn to write & hand off modular code to Eng.
<b>Eng setup &amp; maintenance</b>	↓ Requires CI/CD infra for unit and integration tests, even for one-off models.



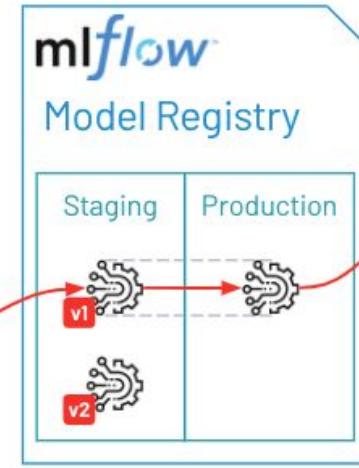
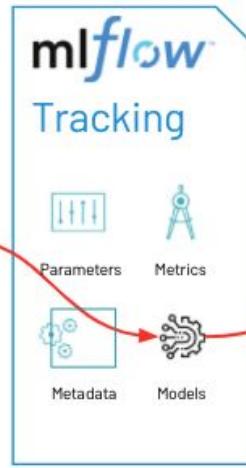
# The Full ML Lifecycle

**Data Scientists** build features.  
**Data Engineers** provide infra for automating featurization.

**Data Scientists** build models and log them to MLflow, which records environment info.

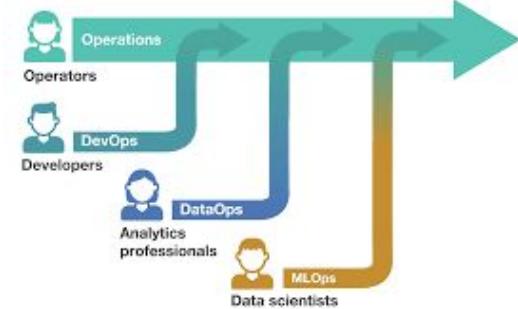
**Data Scientists** move models to Staging.

**Deployment Engineers** manage CI/CD tools which promote models to Production.



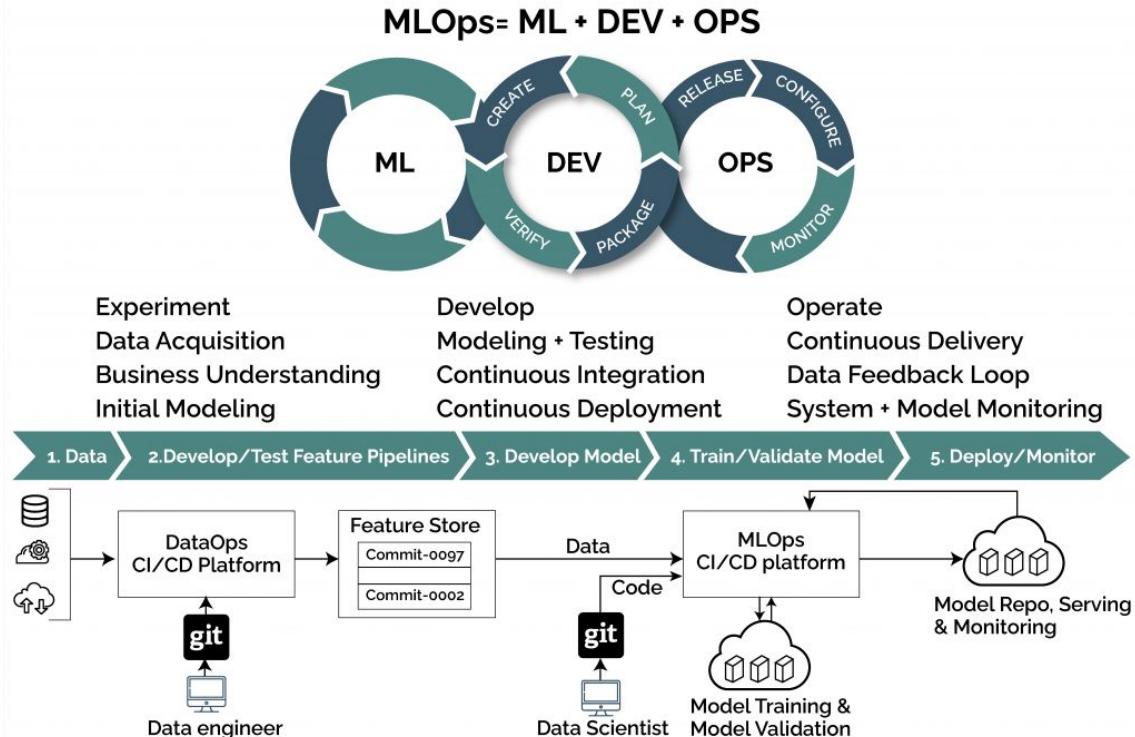
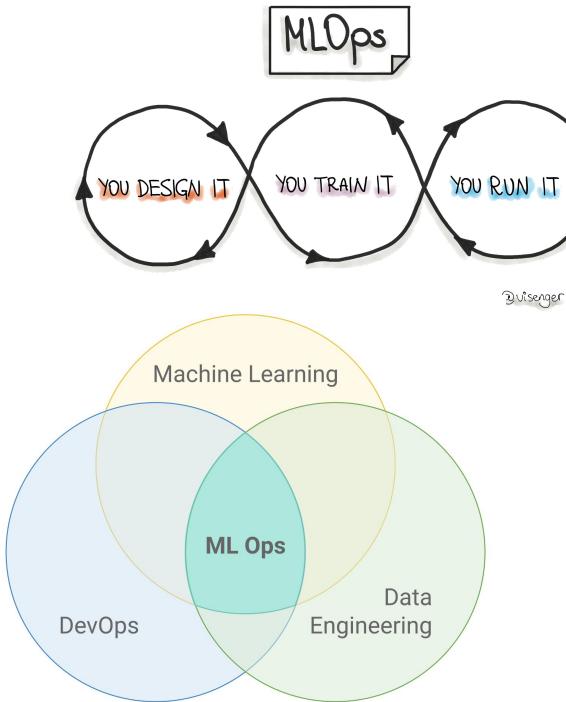
# MLOps

Making ML Systems Systematic, Reliable, and Repeatable



## ModelOps from Model Centric to Data Centric AI

AI System = Code (Model/Algorithm) + Data



# From Big Data to Good Data

- Good data is defined as
  - Unambiguous labels
  - Good cover of important use cases
  - Timely feedback from production data (distribution covers both data drift & concept drift)
  - Sized appropriately

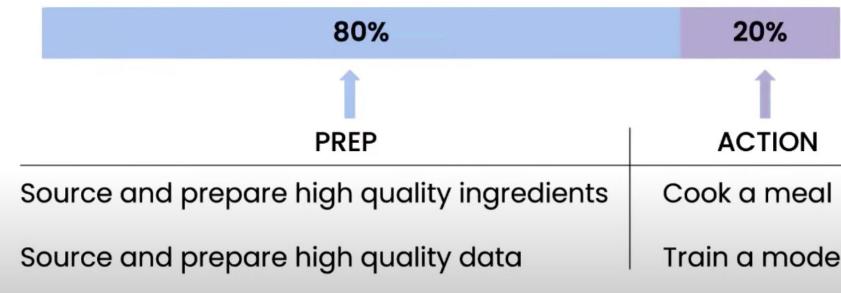
## Model-centric AI

How can you change the model (code) to improve performance?

## Data-centric AI

How can you systematically change your data (inputs x or labels y) to improve performance?

- Data is food for AI
  - Improving the code Vs the data



	Steel defect detection	Solar panel	Surface inspection
Baseline	76.2%	75.68%	85.05%
Model-centric	+0% (76.2%)	+0.04% (75.72%)	+0.00% (85.05%)
Data-centric	+16.9% (93.1%)	+3.06% (78.74%)	+0.4% (85.45%)

# MLOps Progression

## Key Capabilities

### PEOPLE



### DATA ARCH.



### MODEL ARCH.



### PROCESS



### GOVERNANCE



## Beginner

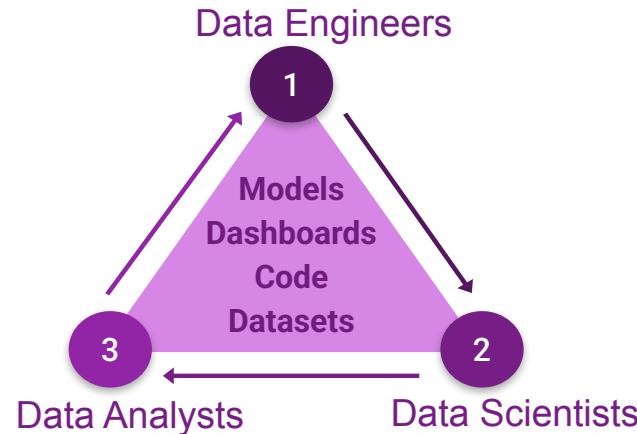
- Use familiar tools
- Increase productivity
- Plan for advanced and more complex workloads

## Intermediate

- Scale data & workloads
- Automate workloads & focus on reproducibility
- Unify data teams & Improve collaboration across all data & AI workloads

## Advanced

- Faster production cycles
- CI/CD & well defined processes
- E2E automation & reproducibility across multiple workloads



# Model Tuning

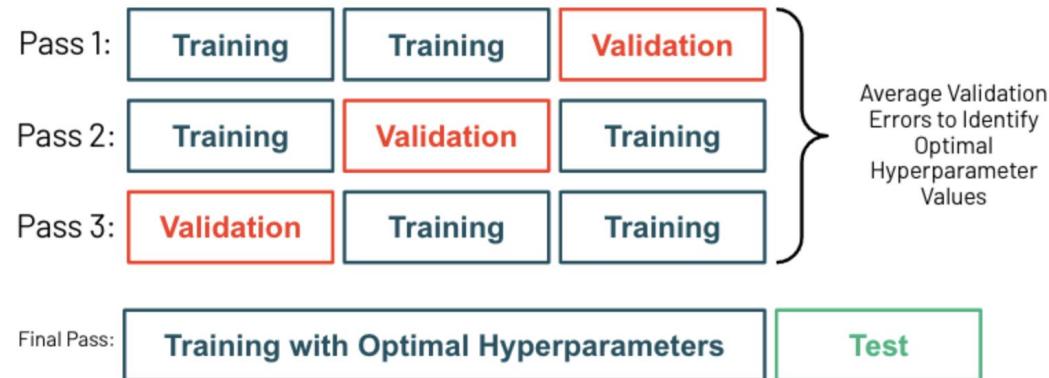
Spark MLlib supports **grid search** for hyperparameter tuning.

A diagram illustrating a grid search. On the left, there is a table with two rows: Tree Depth 5, Number of Trees 2; and Tree Depth 8, Number of Trees 4. A red arrow points from this table to another table on the right with four rows: Tree Depth 5, Number of Trees 2; Tree Depth 5, Number of Trees 4; Tree Depth 8, Number of Trees 2; and Tree Depth 8, Number of Trees 4.

Tree Depth	Number of Trees
5	2
8	4

Tree Depth	Number of Trees
5	2
5	4
8	2
8	4



## Model Selection (aka Hyperparameter Tuning)

- Manual:
  - Grid Search
- Automated:
  - Random Search
  - Bayesian
  - Genetic

## K-fold Cross Validation

Avoid overfit/underfit

Train: Validation: Test

We need a **validation set** to assess model performance within our optimization process

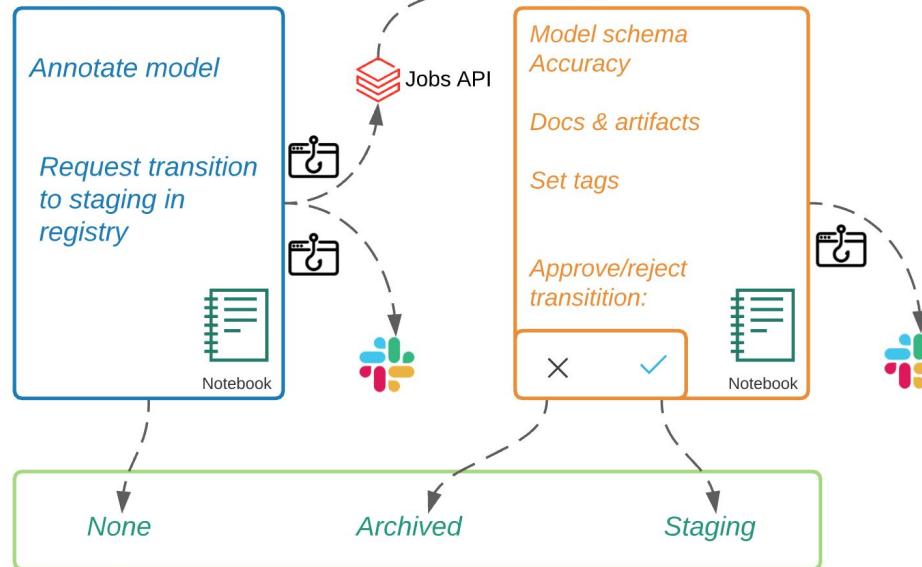
This keeps us from optimizing on our **test** or holdout set.

## Automation with webhooks

Webhook

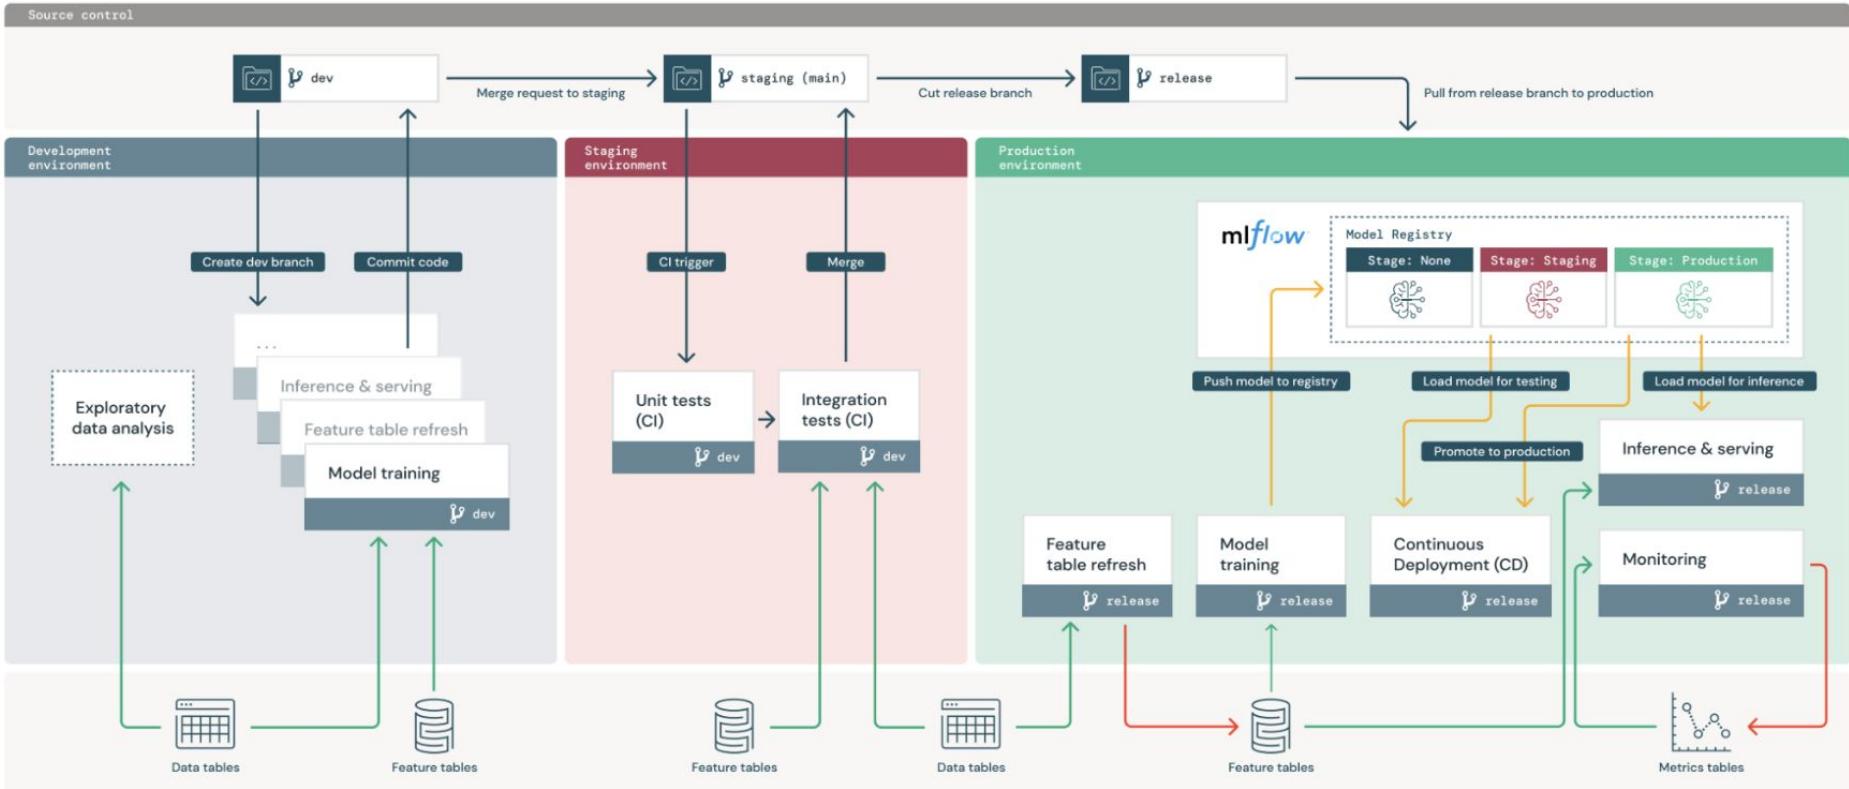
Slack notification

### Promote Best Run to Registry



**mlflow™ Model Registry**

# Overview



# Solution Accelerators

## E2E ML Example

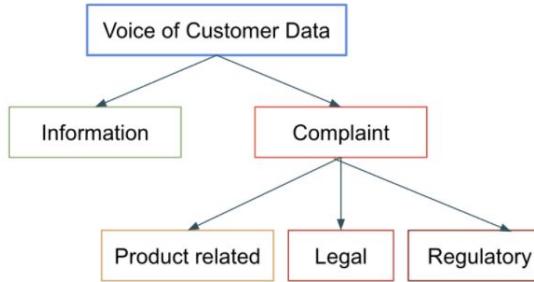
<https://databricks.com/solutions#by-industry>

<https://databricks.com/solutions/accelerators>

The Big Book of Machine Learning Use Case.pdf

<https://databricks.com/p/ebook/big-book-of-machine-learning-use-cases>

# Model Ensembles

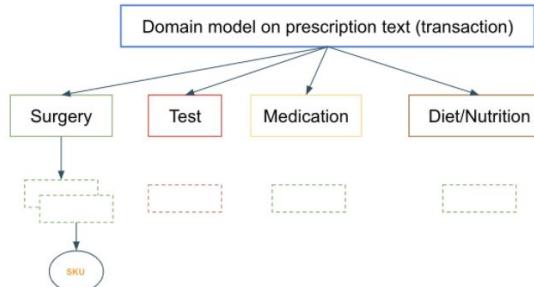


Single Best Model Vs Ensemble of Models  
Champion & Challenger models

## Link

**Multiple models do not necessarily mean an ensemble!**

A layering approach where moderately performant un-correlated models are combined to produce a supermodel that improves accuracy while improving stability and is often a divide and conquer strategy used in large, diverse datasets.



- Fetch the “best” models of each architecture type
- Build a custom pyfunc model class that encapsulates the best models
- Provide a predict function for the ensemble
- **Provide a voting function**
- Package and log the model in MLflow as a custom pyfunc model
- Scoring