

CSCI E-103

Data Engineering for Analytics to Solve Business Challenges

Getting the most value out of your data initiative investments
through a culture of Continuous Improvement

Lecture 13

Eric Gieseke & Ramdas Marali

Harvard Extension, Fall 2025

Agenda

- Final Assignment
- Software Development Lifecycle (SDLC) of a Data Persona
- Creating Robust & Scalable Data Products & Services that are also price-performant
 - Fierce competition in the tech landscape
 - Value added Differentiators
 - Balancing: Capabilities Vs Performance Vs Price
- Automation of Infrastructure & Pipelines(Data/ML) - IaaC & CI/CD
- Observability & Monitoring
- Improving ROI with a Center of Excellence (CoE)
 - Time to Market & Upskilling
 - Capacity Planning & Forecasting
- Data Intelligent Platforms in the age of GenAI
- Lab:
 - Databricks Asset Bundles (DAB)

Final Project - Building the Lakehouse Architecture

Assignment Description

Presentation Agenda

- Problem statement
- Team Introduction (by role)
- Design & Architecture of the Lakehouse Paradigm (Data Architect)
 - Scalable/Quality/Performance/Reliability/Governance
- Data Pipeline - Ingestion/Transformation/Joins/Aggregations (Data Engineer)
- Model - training/Inferencing/life cycle Management (ML Practitioner)
- BI Dashboard - Queries, Dashboard, Notifications (Data Analyst)
 - Connectivity/Scale/Security Concerns
- Summarize Insights
- Next set of planned refinements

Group Presentation Notes

20 slides max
15 min presentation
5 min Q/A

2-3 external judges

Give your team/consulting firm a name


We'll go from Groups 1-N
Folks in other time zones will go first

Operationalizing Pipelines

- Understand Use Case
 - Business Challenges & domain specific nuances
 - Requirements (functional, non-functional) SLAs
- Build for Scale
 - Investment at risk
- Winning the Architecture
 - Build repeatable patterns using frameworks that are config driven
 - Document Best Practices
 - Early Performance Tuning to showcase efficiencies (Project TCO)
 - Infrastructure Efficiencies (storage life cycle, auto termination)
 - Security & Business Continuity
- Observability
 - Monitoring (data & usage)
- Automation & CI/CD
 - Infrastructure as Code (IaC)
 - Testing Automation
- Center Of Excellence (COE)
 - Expertise, best practices, and support

Software Development Lifecycle (SDLC) of Data Personas

Promoting workloads from lower to higher environments



Service Level Agreement (SLA)

Depends on workload type (as soon as possible is not practical)

Batch

Volume

Processing Time

Streaming

TPS

Tx/Events per second

BI

E2E latency

Concurrency

High Availability

Response Time

Availability = (Uptime / Total Time) x 100%

E.g., 99.999% => 5 minutes/year

Disaster Recovery

Recovery Time Objective (RTO)

Recovery Point Objective (RPO)

Price-Performance



- Choose a platform that provides out of the box price-performance
 - yet gives the control to you to choose from
- Industry-standard benchmarks
 - TPC-DS
 - Analytics / Decision Support (BI)
 - How fast and effectively you query the data warehouse
 - [Databricks Sets Official Data Warehousing Performance Record](#)
 - TPC-DI
 - Data Integration / ETL
 - How fast and effectively you build and load the data warehouse
 - [How We Performed ETL on One Billion Records For Under \\$1 With Delta Live Tables](#)

<https://github.com/anhchu/Metimur>

<https://docs.databricks.com/aws/en/sql/tpcds-eval>

Transaction Processing Performance Council (TPC)

Cost Optimizations

- Don't make assumptions, experimentation is cheap
 - Spark is mature and scales almost linearly
 - Performance requirements come in the form of **development**, **ecosystem**, or **business** needs.
- Better code
 - Leverage dataframe APIs over home grown UDFs
 - Cache during ML training
 - Use binary formats like Delta
- Right infrastructure
 - Instance type, cluster sizing
- Governance
 - Cluster policies
 - Tags
- Autoscaling & auto termination
- Spot instances, spot fleet
- Higher Spark versions
 - Better performance & bug fixes apart from enhanced capabilities
- Pool -> Serverless
- Monitoring & Alerting
 - Usage overview for chargeback & control runaway costs
 - Tags for attribution



CI/CD

CI = Continuous Integration

... Developers continually make new changes. Code is auto tested and promoted.

CD = Continuous Delivery / Deployment


... Continually deploy code into production at click of a button / automatically



Development Workflows

Code developed
in IDE, packaged
& deployed as
libraries

Notebooks are
used for
configuration &
orchestration of
calls to libraries



Infrastructure as Code (IaC)

- Infrastructure as Code is an approach to infrastructure automation based on practices from software development
- Core Practices of IaC:
 1. *Define Everything as Code*
 2. *Continuously Test and Deliver All Work in Progress*
 3. *Build Small, Simple Pieces that you can change independently*
- In Databricks, Infrastructure as Code can pertain to both Cloud Infrastructure and Workspace Infrastructure

Automation Options

	Source	Cloud	Type	Infrastructure	Language	Agent	Master	Community	Maturity
Chef	Open	All	Config Mgmt	Mutable	Procedural	Yes	Yes	Large	High
Puppet	Open	All	Config Mgmt	Mutable	Declarative	Yes	Yes	Large	High
Ansible	Open	All	Config Mgmt	Mutable	Procedural	No	No	Huge	Medium
SaltStack	Open	All	Config Mgmt	Mutable	Declarative	Yes	Yes	Large	Medium
CloudFormation	Closed	AWS	Provisioning	Immutable	Declarative	No	No	Small	Medium
Heat	Open	All	Provisioning	Immutable	Declarative	No	No	Small	Low
Terraform	Open	All	Provisioning	Immutable	Declarative	No	No	Huge	Low

- Configuration Management vs Provisioning

- Mutable Infrastructure vs Immutable Infrastructure

- Procedural vs Declarative


- Master vs Masterless

- Agent vs Agentless


- Large Community vs Small Community

- Mature vs Cutting Edge

- Using Multiple Tools Together



Workspace Organization




Balance of Isolation & Management overhead

Lakehouse and Data Mesh


Business units are organized in **Data Domains** and own their respective sources, data and metadata. Using a Lakehouse, they create domain specific insights from data and offer **Data Products** to other domains using **Unity Catalog** a common data catalog. Compliance to organisational rules and industry regulations is ensured via **Federated Governance**.

With a Lakehouse Platform, data domains can be created on different levels:


- One Workspace using clusters to isolate domains
- Using a separate Workspace per data domain
- Data domains being full Lakehouses



General CI/CD Workflow on Databricks (Azure DevOps)




Automation



- [REST APIs](#)
- [Source Control](#)
- [CI/CD server](#)
- Scripting with cloud formation, terraform
 - Deploy workspace
 - Connect data sources
 - Provision users and groups
 - Create clusters and cluster policies
 - Add permissions for users and groups
 - [Blog](#)
 - [Blog](#)
- Clusters
 - Automated Job Cluster
 - Interactive All-purpose Clusters


Fully configured data environments



Terraform Provider


Cloud Agnostic
Code Reuse
Repeatability
Predictable
Supports all REST APIs

- Compute Resources
- Storage
- Security






Repos API and Git-based Jobs

Run production workloads on Databricks following good CI/CD practices



Observability

- System Tables for operational intelligence
 - Monitor your **consumption** and leverage the lakehouse AI capabilities to forecast your future usage, triggering alerts when billing goes above your criterias
 - Monitor **accesses** to your data assets
 - Monitor and understand your **platform usage**
 - [Db demos](#)




Data teams have SLAs on delivery not quality.

- **Reactive issue management** : data consumers encounter issues with data before the data team
- **Bottlenecked operations** : data consumers don't have self-serve experiences and must go through data team to add or use data
- **Long Root-Cause-Analysis** : it can take weeks to debug quality issues and rollback changes

Data Quality Monitoring Features

Anomaly Detection

- Enabled at the **schema level**
- ✨ Intelligently detects data quality anomalies for all tables
- Meant for: Scalable quality monitoring and actionable insights




Data profiling

- Enabled at the **table level, f.k.a Lakehouse Monitoring**
- 📊 Provides a table profile for your most important tables
- Meant for: DIY table monitoring and quick summary statistics

AI-Powered

Intelligent and built on your data

- **Freshness:** tracks how recently data was updated, flagging tables as stale if updates are delayed.
- **Completeness:** checks if row volume in the last 24 hours falls below expected levels.
- **Segmentation:** View quality metrics by slice (e.g., freshness of vendor_id="walmart")



Bespoke rules per table

Simple and scalable to all tables

Holistic view across all tables

Search data, notebooks, recents, and more... ⌘ + P

PrPreview Pr-1060831-7d5d2 E2 Dogfood K

Warehouse_picker

Data quality monitoring

selected_catalog_name ▾ All schemas ▾

Data quality ⓘ 95.87% healthy As of 2 hours ago

2756 tables healthy 115 tables unhealthy

Tables monitored ⓘ 2,871 Configure

36% of all tables

Quality incidents

Important 21 All issues 115


Schema	Table	Reason	Since	Impact	Root cause	Related
schema_n1	marketing_campaign...	Stale	2 hr ago	High	campaign_etl_daily job failed	Fixes 4 incidents
schema_n2	finance_transactions...	Incomplete	1 hr ago	High	finance_ingest_job job failed	
schema_n3	compliance_audit_logs	Stale	45 min ago	High	audit_log_etl_job job failed	Fixes 2 incidents
schema_n4	product_line_info	Stale	37 min ago	High	campaign_etl_daily job failed	
schema_n5	customer_transactions	Stale	1 hr ago	High	campaign_etl_daily job failed	
schema_n6	store_info	Incomplete	12 min ago	Medium	campaign_etl_daily job failed	
schema_n7	sales_pipeline_gold	Incomplete	15 min ago	Medium	audit_log_etl_job job failed	
schema_n8	finance_transactions...	Stale	1 hr ago	Medium	-	
schema_n9	product_line_info	Stale	2 hr ago	Medium	-	
schema_n10	compliance_audit_logs	Stale	4 hr ago	Medium	-	

Load more ▾

Incidents sorted by downstream impact

Links to the root cause in the upstream job


Whether data comes from
data ingestion or model
outputs,



data manifests as **tables**.


Monitoring a table in the Lakehouse

How does it work?



Data Profiling*

Automated profiling of your most important tables





DQX - Data Quality Framework

Provided by [Databricks Labs](#)

DQX is a data quality framework for Apache Spark that enables you to define, monitor, and address data quality issues in your Python-based data pipelines.

[Motivation](#)

[Installation](#)

[User guide](#)

[Demos](#)

[Reference](#)